

Essential R: Data Analysis, Visualization, and Modeling

Anju Sambasiavn

2024-08-03

Contents

1	Part-1	3
1.1	getwd(): Get info of current working directory.	3
1.2	Types of variable	3
1.3	Variable Assignment	4
1.4	Checking Variable Types	4
1.5	Operators in R	5
1.6	Sequence Control	8
2	Part-2	12
2.1	Matrix	12
2.2	Creating Matrices	12
2.3	Filling by Row	13
2.4	Accessing Matrix Elements	13
2.5	Matrix operations	14
2.6	Matrix Functions	16
2.7	Column Binding (<code>cbind</code>)	17
2.8	Naming and Accessing Elements	17
2.9	Vector	19
2.10	Key Characteristics of Vectors in R	19
2.11	Different Methods for Vector Creation	19
2.12	Vector Indexing and Subsetting	21
2.13	Vector Operations	23
2.14	Application Level	24
2.15	Questions	29
3	Part-3	35
3.1	Read the Data from a CSV File	35
3.2	Explore and Understand the Data	35
3.3	Subsetting the Data Frame	36
3.4	Accessing Columns Using <code>\$</code>	37
3.5	Create an Additional Attribute (Add a Column)	38
3.6	Filter the Data	39
3.7	Creating a New Attribute Based on Conditions	40
3.8	Quick Plotting with <code>qplot</code>	41
4	Part-4	46
4.1	Reading the CSV Files	46
4.2	Exploring the Dataframes	46
4.3	Load Data	47
4.4	Data Preprocessing	47
4.5	Data Visualization with <code>ggplot2</code>	48

4.6	Settings vs. Mappings	56
4.7	Geometric and Statistical Plots	60
4.8	Exercises and Comparisons	67
4.9	Loading the Data and Renaming Columns	78
4.10	Creating a Histogram with a Specific Color	79
4.11	Focusing on a Specific Range Using <code>coord_cartesian()</code>	80
4.12	Creating a Scatter Plot and Applying Axis Limits	82
4.13	Summary	84
4.14	Loading Data and Preparing the Base Plot	84
4.15	Facets of Histograms	84
4.16	Facets of Scatterplots	88
4.17	Customizing Plot Themes	93
4.18	Moving the Legend Inside the Plot Area	95
4.19	Adding a Title to the Plot	96
5	Part-5	100
5.1	Loading Data and Summarizing	100
5.2	Basic Statistical Functions for Price	100
5.3	Correlation Between Miles and Price for Model X	101
5.4	Pearson's Product-Moment Correlation Interpretation	101
5.5	Scatter Plot of Miles vs. Price	102
5.6	Simple Linear Regression (SLR)	103
5.7	Linear Regression Output Explained	104
5.8	Multiple Linear Regression Model	107
5.9	Code for Splitting Data and Fitting Linear Regression:	108
5.10	Code for without Splitting Data and Fitting Linear Regression:	108
5.11	Linear Regression Model Interpretation	109
5.12	Interpretation of the Residuals Histogram	111
5.13	Logistic Regression	113
5.14	Loading Libraries and Data	114
5.15	Splitting the Data	114
5.16	Data Wrangling and Model Fitting	114
5.17	Interpretation:	115
5.18	Key Points About Logistic Regression	116
5.19	Making Predictions and Evaluating the Model	117
5.20	Confusion matrix:	118
6	Part-6	121
6.1	Neural Network (NN)	121
6.2	Creating a Training Dataset	121
6.3	Fitting the Neural Network	121
6.4	Predicting New Data	121
6.5	Exercise - Neural Network Model on Dataset	122

1 Part-1

- Programming Language: Designed for statistical computing and data analysis.
- Free Software: Open-source and free to use.
- Statistical Tools: Includes a wide range of statistical and graphical techniques.
- Data Visualization: Excellent for creating plots and charts.
- Extensible: Allows users to add custom functions and packages.
- Popular in Academia: Widely used in research and education.
- Community Support: Strong user community and numerous online resources.

1.1 getwd(): Get info of current working directory.

```
getwd()
```

```
[1] "C:/Desktop/3MonthWorks/BusinessIntelligence/R-Essentials/R-Essentials"
```

1.2 Types of variable

- Characters
- Numeric(real number)
- Logical
- Integer
- Factor
- Complex

1.2.1 Characters:

- Character variables store text data.

```
name <- "Anju"  
city <- "Christchurch"  
name
```

```
[1] "Anju"
```

```
city
```

```
[1] "Christchurch"
```

1.2.2 Numeric:

- Numeric variables store numerical data such as integers or decimals.

```
age <- 30  
temperature <- 25.5  
age
```

```
[1] 30
```

```
temperature
```

```
[1] 25.5
```

1.2.3 Logical:

- Logical variables store boolean values, which can be either TRUE or FALSE.

```
is_student <- TRUE
has_car <- FALSE
is_student
```

```
[1] TRUE
```

```
has_car
```

```
[1] FALSE
```

1.2.4 Integer:

- Integer variables store whole numbers.

```
count <- 10L # L suffix indicates integer type
count
```

```
[1] 10
```

1.2.5 Factor:

- Factor variables are used to represent categorical data with levels.

```
gender <- factor(c("Male", "Female", "Male", "Female"))
gender
```

```
[1] Male   Female Male   Female
```

```
Levels: Female Male
```

1.2.6 Complex:

- Complex variables store complex numbers with real and imaginary parts.

```
z <- 3 + 2i
z
```

```
[1] 3+2i
```

1.3 Variable Assignment

- In R, variables are assigned using the <- operator (though “=” can also be used).

```
z <- 5
z
```

```
[1] 5
```

1.4 Checking Variable Types

- We can check the type of a variable using the class() function.

```
class(age)
```

```
[1] "numeric"
```

```
class(name)
```

```
[1] "character"
```

```
class(is_student)
```

```
[1] "logical"
```

1.5 Operators in R

1. Arithmetic operators
2. Relational operators
3. Logical operators
4. Assignment operators
5. Miscellaneous operators

1.5.1 Arithmetic Operator:

- Arithmetic operators are used to perform basic mathematical operations.

```
a <- 5  
b <- 3
```

```
add <- a + b  
cat("Addition: ", add, "\n")
```

```
Addition: 8
```

```
sub <- a - b  
cat("Subtraction: ", sub, "\n")
```

```
Subtraction: 2
```

```
mult <- a * b  
cat("Multiplication: ", mult, "\n")
```

```
Multiplication: 15
```

```
div <- a / b  
cat("Division: ", div, "\n")
```

```
Division: 1.666667
```

```
exp <- a ^ b  
cat("Exponentiation: ", exp, "\n")
```

```
Exponentiation: 125
```

```
mod <- a %% b  
cat("Modulus: ", mod, "\n")
```

```
Modulus: 2
```

```
intdiv <- a %/% b  
cat("Integer Division: ", intdiv, "\n")
```

```
Integer Division: 1
```

1.5.2 Relational Operators:

- Relational operators compare values and return logical values (TRUE or FALSE).

```
a <- 5
b <- 3
```

```
equal_result <- a == b
cat("Equal to (a == b):", equal_result, "\n")
```

Equal to (a == b): FALSE

```
not_equal_result <- a != b
cat("Not equal to (a != b):", not_equal_result, "\n")
```

Not equal to (a != b): TRUE

```
greater_than_result <- a > b
cat("Greater than (a > b):", greater_than_result, "\n")
```

Greater than (a > b): TRUE

```
less_than_result <- a < b
cat("Less than (a < b):", less_than_result, "\n")
```

Less than (a < b): FALSE

```
greater_than_or_equal_result <- a >= b
cat("Greater than or equal to (a >= b):", greater_than_or_equal_result, "\n")
```

Greater than or equal to (a >= b): TRUE

```
less_than_or_equal_result <- a <= b
cat("Less than or equal to (a <= b):", less_than_or_equal_result, "\n")
```

Less than or equal to (a <= b): FALSE

1.5.3 Logical Operators:

- Logical operators are used to combine multiple conditions.

```
x <- TRUE
y <- FALSE
```

```
and_result <- x & y
cat("Logical AND (x & y):", and_result, "\n")
```

Logical AND (x & y): FALSE

```
or_result <- x | y
cat("Logical OR (x | y):", or_result, "\n")
```

Logical OR (x | y): TRUE

```
not_result <- !x
cat("Logical NOT (!x):", not_result, "\n")
```

Logical NOT (!x): FALSE

```
xor_result <- xor(x, y)
cat("Logical XOR (xor(x, y)):", xor_result, "\n")
```

Logical XOR (xor(x, y)): TRUE

```
and_multiple_result <- (x & !y) & (x | y)
cat("Logical AND with multiple conditions:", and_multiple_result, "\n")
```

Logical AND with multiple conditions: TRUE

1.5.4 Assignment Operators:

- Assignment operators are used to assign values to variables.

```
a <- 10
b <- 5

a <- b
cat("After b assigned to a: a =", a, "\n")
```

After b assigned to a: a = 5

```
my_function <- function(x = 5) {
  return(x)
}

print(my_function())
```

[1] 5

```
print(my_function(10))
```

[1] 10

1.5.5 Miscellaneous Operators

```
df <- data.frame(
  ID = 1:5,
  Name = c("Alice", "Bob", "Charlie", "David", "Eve"),
  Age = c(25, 30, 35, 40, 45)
)

# 1. Colon Operator (:)
# Creates a sequence of numbers from 1 to 10
sequence <- 1:10
cat("Sequence created using colon operator (1:10):\n")
```

Sequence created using colon operator (1:10):

```
print(sequence)
```

[1] 1 2 3 4 5 6 7 8 9 10

```
# 2. Membership (%in%)
# Check if elements are in a vector
vector <- c(1, 3, 5, 7, 9)
membership_check <- c(2, 3, 4) %in% vector
cat("Membership check (c(2, 3, 4) %in% vector):\n")
```

Membership check (c(2, 3, 4) %in% vector):

```
print(membership_check)
```

[1] FALSE TRUE FALSE

```
# 3. Concatenation (c())
# Combine elements into a vector
combined_vector <- c(1, 2, 3, 4, 5)
cat("Concatenated vector (c(1, 2, 3, 4, 5)):\n")
```

Concatenated vector (c(1, 2, 3, 4, 5)):

```
print(combined_vector)
```

```
[1] 1 2 3 4 5
```

```
# 4. Subset using $ (extract a column from a data frame)
name_column <- df$Name
cat("Extracted Name column using $:\n")
```

Extracted Name column using \$:

```
print(name_column)
```

```
[1] "Alice" "Bob" "Charlie" "David" "Eve"
```

```
# 5. Subset using [ (extract rows and columns from a data frame)
subset_rows <- df[1:3, ]
cat("Subset of first 3 rows using [:\n")
```

Subset of first 3 rows using [:

```
print(subset_rows)
```

```
ID    Name Age
1 1    Alice 25
2 2      Bob 30
3 3  Charlie 35
```

```
# 6. Subset using [[ (extract a single element or list element)
age_column <- df[["Age"]]
cat("Extracted Age column using [[:\n")
```

Extracted Age column using [[:

```
print(age_column)
```

```
[1] 25 30 35 40 45
```

1.6 Sequence Control

1. Conditional statements
 - a. if
 - b. if...else
 - c. if...else if...else
2. Loops
 - a. for
 - b. while
 - c. repeat
3. Control statements

- a. break
- b. next

1.6.1 Conditional Statements

1.6.1.1 if Statement:

- It evaluates a condition and executes a block of code if the condition is TRUE.

```
x <- 10
if(x > 5){
  print("x is greater than 5")
}
```

```
[1] "x is greater than 5"
```

1.6.1.2 if...else:

- It allows you to execute one block of code if the condition is TRUE and another block if it is FALSE.

```
x <- 3
if (x > 5) {
  print("x is greater than 5")
} else {
  print("x is not greater than 5")
}
```

```
[1] "x is not greater than 5"
```

1.6.1.3 if...else if...else Statement:

- It allows you to check multiple conditions sequentially.

```
x <- 7
if (x < 5){
  print("x is less than 5")
} else if(x >= 5 & x < 10){
  print("x is between 5 and 9")
} else {
  print("x is 10 or greater")
}
```

```
[1] "x is between 5 and 9"
```

1.6.2 Loops

- Loops are used to repeat a block of code multiple times until a specified condition is met.

1.6.2.1 for:

- It iterates over a sequence (e.g., a vector or a sequence of numbers) and executes a block of code for each element.

```
for (i in 1:5){
  print(i)
}
```

```
[1] 1
[1] 2
```

```
[1] 3
[1] 4
[1] 5
```

1.6.2.2 while:

- It repeats a block of code as long as a specified condition is TRUE.

```
x <- 1
while(x <= 5){
  print(x)
  x <- x + 1
}
```

```
[1] 1
[1] 2
[1] 3
[1] 4
[1] 5
```

1.6.2.3 repeat:

- It repeatedly executes a block of code until a break statement is encountered.

```
x <- 1
repeat {
  print(x)
  x <- x + 1
  if (x > 5){
    # Exit the loop when x > 5
    break
  }
}
```

```
[1] 1
[1] 2
[1] 3
[1] 4
[1] 5
```

1.6.3 Control Statements

1.6.3.1 break:

- This is used to exit a loop prematurely.

```
for(i in 1:10){
  if (i>5) {
    break # Exit the loop when i >5
  }
  print(i)
}
```

```
[1] 1
[1] 2
[1] 3
[1] 4
[1] 5
```

1.6.3.2 next:

- It skips the current iteration of a loop and continues with the next iteration.

```
for(i in 1:5) {  
  if(i == 3) {  
    next # Skip iteration when i=3  
  }  
  print(i)  
}
```

```
[1] 1  
[1] 2  
[1] 4  
[1] 5
```

1.6.3.3 return:

- It is used to exit a function and return a value.

```
my_function <- function(x) {  
  if(x < 0){  
    return("Input is negative")  
  } else{  
    return("Input is positive")  
  }  
}
```

```
result <- my_function(-5)  
print(result)
```

```
[1] "Input is negative"
```

2 Part-2

2.1 Matrix

- A matrix is a two-dimensional array that holds elements arranged in rows and columns.
 - Matrices are essentially a collection of vectors arranged in a grid.
1. Dimension: Two-dimensional (rows and columns).
 2. Type: Homogeneous (all elements must be of the same type).
 3. Indexing: Accessed by two indices (row and column).

2.2 Creating Matrices

- Create a matrix using the `matrix()` function.
- The `matrix()` function takes a vector of elements and organizes them into a matrix of specified dimensions.

```
# Create a 3x3 matrix with numbers 1 to 9  
# By default, matrices are filled column-wise.  
m <- matrix(1:9, nrow = 3, ncol = 3)  
print(m)
```

```
      [,1] [,2] [,3]  
[1,]     1     4     7  
[2,]     2     5     8  
[3,]     3     6     9
```

```
# Column wise  
my_data <- 1:20  
A <- matrix(my_data, 4, 5)
```

```
# Creates a 4x5 matrix with numbers 1 to 20. Filled column-wise by default.  
my.data <- 1:20  
A <- matrix(my.data, 4, 5)  
print(A)
```

```
      [,1] [,2] [,3] [,4] [,5]  
[1,]     1     5     9    13    17  
[2,]     2     6    10    14    18  
[3,]     3     7    11    15    19  
[4,]     4     8    12    16    20
```

2.2.1 Total Elements = row*col

```
# Total elements = 3 * 4 = 12  
data <- 1:12  
matrix_data <- matrix(data, nrow = 3, ncol = 4)  
print(matrix_data)
```

```
      [,1] [,2] [,3] [,4]  
[1,]     1     4     7    10  
[2,]     2     5     8     9  
[3,]     3     6     9    12
```

2.2.2 Recycle the data

```
# Total elements = 3 * 4 = 12, but data has only 8 elements  
# This will recycle the data.  
data <- 1:8  
matrix_data <- matrix(data, nrow = 3, ncol = 4)
```

Warning in matrix(data, nrow = 3, ncol = 4): data length [8] is not a sub-multiple or multiple of the number of rows [3]

```
print(matrix_data)
```

```
      [,1] [,2] [,3] [,4]  
[1,]    1    4    7    2  
[2,]    2    5    8    3  
[3,]    3    6    1    4
```

2.3 Filling by Row

```
#Create a 3x3 matrix filled by row  
m_byrow <- matrix(1:9, nrow = 3, byrow = TRUE)  
print(m_byrow)
```

```
      [,1] [,2] [,3]  
[1,]    1    2    3  
[2,]    4    5    6  
[3,]    7    8    9
```

2.4 Accessing Matrix Elements

```
# Single Element: Access a specific element using row and column indices.  
m <- matrix(2:13, nrow = 4, ncol = 3)  
print(m)
```

```
      [,1] [,2] [,3]  
[1,]    2    6   10  
[2,]    3    7   11  
[3,]    4    8   12  
[4,]    5    9   13
```

```
element <- m[2, 3]  
print(element)
```

```
[1] 11
```

```
# Row wise  
A <- matrix(my_data, 4, 5, byrow = TRUE)
```

```
# Accesses the element in the 2nd row, 3rd column: 8  
a1 <- A[2, 3]  
a1
```

```
[1] 8
```

```
# Accesses the entire 2nd row: 6 7 8 9 10  
a2 <- A[2, ]
```

```

a2

[1] 6 7 8 9 10
# Accesses the entire 3rd column: 3 8 13 18
a3 <- A[, 3]
a3

[1] 3 8 13 18

```

2.5 Matrix operations

Matrix operations are fundamental in linear algebra and data manipulation, especially in programming and data science.

- Matrix Addition
- Element-wise Multiplication
- Matrix Multiplication
- Transpose

2.5.1 Matrix Addition:

- It involves adding corresponding elements of two matrices of the same dimensions.

```

# Matrix Addition
m1 <- matrix(1:8, nrow = 4, byrow = TRUE)
m2 <- matrix(8:15, nrow = 4)
result_add <- m1 + m2
print(result_add)

```

```

      [,1] [,2]
[1,]     9    14
[2,]    12    17
[3,]    15    20
[4,]    18    23

```

Explanation:

- m1 and m2 are matrices of the same size (4x2 in this case).
- Addition is performed element-wise: (m1[1, 1] + m2[1, 1]) for the element in the first row, first column, and so on.

2.5.2 Element-wise Multiplication

Element-wise Multiplication multiplies corresponding elements of two matrices of the same dimensions.

```

# Element-wise Multiplication
result_mult <- m1 * m2
print(result_mult)

```

```

      [,1] [,2]
[1,]     8    24
[2,]    27    52
[3,]    50    84
[4,]    77   120

```

Explanation:

- `m1 * m2` multiplies each element of `m1` by the corresponding element in `m2`.

2.5.3 Matrix Multiplication

Matrix Multiplication (also known as matrix product) involves a more complex operation than element-wise multiplication. It requires that the number of columns in the first matrix be equal to the number of rows in the second matrix.

```
# Matrix Multiplication
m3 <- matrix(1:8, nrow = 4, byrow = TRUE)
m4 <- matrix(8:15, nrow = 2, ncol = 4)
result_mat_mul <- m3 %*% m4
print(result_mat_mul)
```

```
      [,1] [,2] [,3] [,4]
[1,]   26   32   38   44
[2,]   60   74   88  102
[3,]   94  116  138  160
[4,]  128  158  188  218
```

Explanation:

- `%*%` performs matrix multiplication.
- For two matrices A (dimensions $m \times n$) and B (dimensions $n \times p$), the resulting matrix C will have dimensions $m \times p$.
- `c[i, j]` is computed as the sum of the products of the elements of the i -th row of A and the j -th column of B.

2.5.4 Transpose

Transpose of a matrix swaps its rows with columns.

```
# Transpose
m_transpose <- t(result_mat_mul)
print(m_transpose)
```

```
      [,1] [,2] [,3] [,4]
[1,]   26   60   94  128
[2,]   32   74  116  158
[3,]   38   88  138  188
[4,]   44  102  160  218
```

Explanation:

- `t()` function swaps rows and columns of the matrix.

2.5.5 Summary of Matrix Operations:

1. **Addition (+):** Adds corresponding elements of two matrices of the same size.
2. **Element-wise Multiplication (*):** Multiplies corresponding elements of two matrices of the same size.
3. **Matrix Multiplication (%*%):** Multiplies two matrices where the number of columns in the first matrix matches the number of rows in the second matrix.
4. **Transpose (t()):** Swaps rows and columns of a matrix.

2.6 Matrix Functions

2.6.1 dim()

Purpose: Returns the dimensions of a matrix.

```
dimensions <- dim(m_transpose)
print(dimensions)
```

```
[1] 4 4
```

Explanation:

- dim() function returns a vector with two elements: The number of rows and the number of columns of the matrix.
- Example output could be (4, 4) if m_transpose is a 4x4 matrix.

2.6.2 sum()

Purpose: Calculates the sum of all elements in the matrix.

```
total_sum <- sum(m_transpose)
print(total_sum)
```

```
[1] 1664
```

Explanation:

- sum() computes the total sum of all elements in the matrix.
- If m_transpose is a matrix with elements 1, 2, 3, 4, then sum(m_transpose) would be 10.

2.6.3 rowSums()

Purpose: Calculates the sum of elements for each row in the matrix.

```
r_sums <- rowSums(m_transpose)
print(r_sums)
```

```
[1] 308 380 452 524
```

Explanation:

- rowSums() returns a vector where each element is the sum of the elements in the corresponding row of the matrix.

2.6.4 colSums()

Purpose: Calculates the sum of elements for each column in the matrix.

```
c_sums <- colSums(m_transpose)
print(c_sums)
```

```
[1] 140 324 508 692
```

Explanation:

- colSums() returns a vector where each element is the sum of the elements in the corresponding column of the matrix.

Purpose: Combines multiple vectors or matrices into a single matrix by stacking them as rows.


```

r1 <- c("hello", "world", "today")
r2 <- c("mon", "tue", "wed")
r3 <- c(3, 4, 5) # Mixed data types

m1 <- rbind(r1, r2, r3)
print(m1)

```

```

      [,1]      [,2]      [,3]
r1 "hello" "world" "today"
r2 "mon"   "tue"   "wed"
r3 "3"     "4"     "5"

```

Explanation:

- `rbind()` function combines the vectors `r1`, `r2`, and `r3` into a matrix, stacking them as rows.
- Note: Since `r1` and `r2` are character vectors and `r3` is numeric, all data is coerced to character type to accommodate mixed types.

2.7 Column Binding (`cbind`)

Purpose: Combines multiple vectors or matrices into a single matrix by stacking them as columns.

```

c1 <- 1:5
c2 <- -2:-6
m2 <- cbind(c1, c2)
print(m2)

```

```

      c1 c2
[1,]  1 -2
[2,]  2 -3
[3,]  3 -4
[4,]  4 -5
[5,]  5 -6

```

Explanation:

- `cbind()` function combines `c1` and `c2` into a matrix, placing them as columns.

2.8 Naming and Accessing Elements

2.8.1 Naming Vectors

Purpose: Assign names to elements in a vector and access them by name.

```

v1 <- 5:9
names(v1) <- c("a", "b", "c", "d", "e")
print(v1)

```

```

a b c d e
5 6 7 8 9

```

```

v1["d"]

```

```

d
8

```

Explanation:

- `names(v1) <- c("a", "b", "c", "d", "e")` assigns names to the elements of the vector `v1`.

- We can access elements using these names.
- For instance, `v1["d"]` retrieves the value associated with the name “d”.

2.8.2 Removing Names from a Vector

```
names(v1) <- NULL
```

Explanation:

- `names(v1) <- NULL` removes names from the vector `v1`.

2.8.3 Matrix Creation and Naming

Purpose: Create a matrix and assign row and column names.

```
# A vector with 3 elements: "a", "B", and "hello"
v1 <- c("a", "B", "hello")

# Repeats the elements of v1 three times:
# "a", "B", "hello", "a", "B", "hello", "a", "B", "hello"
v2 <- rep(v1, 3)

# Repeats each element of v1 three times in sequence:
# "a", "a", "a", "B", "B", "B", "hello", "hello", "hello"
v3 <- rep(v1, each = 3)

mat <- matrix(v3, nrow = 3, ncol = 3)
rownames(mat) <- c("how", "are", "you")
colnames(mat) <- c("apple", "banana", "kiwi")
print(mat)
```

```
      apple banana kiwi
how "a"      "B"      "hello"
are "a"      "B"      "hello"
you "a"      "B"      "hello"
```

Explanation:

- `matrix(v3, nrow = 3, ncol = 3)` creates a 3x3 matrix using `v3`.
- `rownames(mat)` and `colnames(mat)` assign names to rows and columns respectively.

2.8.4 Accessing Elements Using Names:

```
mat["how", "kiwi"]
```

```
[1] "hello"
```

Explanation:

- Access elements by specifying row and column names.
- For instance, `mat["how", "kiwi"]` retrieves the element at the intersection of the row “how” and column “kiwi”.

2.8.5 Removing Row and Column Names

```
rownames(mat) <- NULL  
colnames(mat) <- NULL
```

Explanation:

- `rownames(mat) <- NULL` and `colnames(mat) <- NULL`
- remove row and column names from the matrix.

2.8.6 Summary

- **Matrix Functions:** `dim()`, `sum()`, `rowSums()`, and `colSums()` help in analyzing matrix dimensions and summarizing data.
- **Row and Column Binding:** `rbind()` and `cbind()` are used to combine matrices or vectors by rows or columns.
- **Naming and Accessing Elements:** Assign and access names in vectors and matrices to make data manipulation more intuitive.

2.9 Vector

In R, a vector is a fundamental data structure used to store elements of the same type. Vectors are essential for handling and manipulating data in R because they allow for efficient and convenient operations on data collections.

2.10 Key Characteristics of Vectors in R

- **Homogeneous Elements:** All elements in a vector must be of the same data type. For example, a numeric vector can only contain numbers, a character vector can only contain strings, and so forth.
- **One-Dimensional:** Vectors are one-dimensional arrays, meaning they only have a single axis. They can be thought of as a list or sequence of elements.
- **Indexed:** Elements in a vector are accessed via indices, which start from 1 in R. For example, `v[1]` accesses the first element of the vector `v`.

2.11 Different Methods for Vector Creation

2.11.1 Using `c()` Function

Purpose: The `c()` function combines elements into a vector.

a. Numeric Vector:

```
nums <- c(1, 2, 3, 4, 5)  
nums
```

```
[1] 1 2 3 4 5
```

Explanation:

- Creates a numeric vector with elements 1 through 5.

b. Character Vector:

```
chars <- c("apple", "banana", "orange")  
chars
```

```
[1] "apple" "banana" "orange"
```

Explanation:

- Creates a character vector with three fruit names.

c. Logical Vector:

```
logic <- c(TRUE, FALSE, TRUE)
logic
```

```
[1] TRUE FALSE TRUE
```

Explanation:

- Creates a logical vector with boolean values.

d. Mixed Type:

```
vec <- c("a", 2, 3, "b")
vec
```

```
[1] "a" "2" "3" "b"
```

Explanation:

- All elements are coerced to character type: c("a", "2", "3", "b").

2.11.2 Other Ways to Create Vectors: seq() and rep()**a. Sequence of Numbers:**

```
vec0 <- 6:12
num_seq <- seq(from = 1, to = 10, by = 2)
vec1 <- seq(1, 15)
vec2 <- seq(1, 15, 2)
num_seq
```

```
[1] 1 3 5 7 9
```

Explanation:

- 6:12 creates a sequence from 6 to 12.
- seq(from = 1, to = 10, by = 2) generates a sequence from 1 to 10 with a step of 2: 1, 3, 5, 7, 9.
- seq(1, 15) generates a sequence from 1 to 15 with a default step of 1.
- seq(1, 15, 2) generates a sequence from 1 to 15 with a step of 2.

b. Repeating Elements:

```
nums_rep <- rep(1:3, times = 2)
vec3 <- rep(2, 5)
vec4 <- rep("hello", 3)
vec4
```

```
[1] "hello" "hello" "hello"
```

Explanation:

- rep(1:3, times = 2) repeats the sequence 1, 2, 3 twice: 1, 2, 3, 1, 2, 3.
- rep(2, 5) repeats the number 2 five times: 2, 2, 2, 2, 2.
- rep("hello", 3) repeats the string "hello" three times: "hello", "hello", "hello".

c. Mixed Elements:

```
v2 <- c("h", "ell", "o")
v2 <- c("h", "ell", "o", 7)
v2
```

```
[1] "h"    "ell"  "o"    "7"
```

Explanation:

- v2 initially contains “h”, “ell”, “o”.
- After adding 7, all elements are coerced to character: c(“h”, “ell”, “o”, “7”).

```
vec5 <- c(7, 120)
vec6 <- rep(vec5, 2)
vec5
```

```
[1] 7 120
```

```
vec6
```

```
[1] 7 120 7 120
```

Explanation:

- vec5 is a vector with elements 7 and 120.
- rep(vec5, 2) repeats the vec5 vector twice: 7, 120, 7, 120.

2.12 Vector Indexing and Subsetting

2.12.1 Indexing with [] Bracket:

Explanation:

- Access or remove elements using indexing.
- Negative indices exclude specified elements.

```
w <- c(2, 3, 4, 5, 6, 7, 81, 21)
# First element: 2
w[1]
```

```
[1] 2
```

```
# Second element: 3
w[2]
```

```
[1] 3
```

```
# Fifth element: 6
w[5]
```

```
[1] 6
```

```
# All elements except the first one: 3, 4, 5, 6, 7, 81, 21
w[-1]
```

```
[1] 3 4 5 6 7 81 21
```

```
# All elements except the third one: 2, 3, 5, 6, 7, 81, 21
w[-3]
```

```
[1] 2 3 5 6 7 81 21
```

```
# Elements from the first to the third: 2, 3, 4
w[1:3]
```

```
[1] 2 3 4
```

```
# Elements from the fifth to the seventh: 6, 7, 81
w[5:7]
```

```
[1] 6 7 81
```

```
# All elements except the first to third: 5, 6, 7, 81, 21
w[-1:-3]
```

```
[1] 5 6 7 81 21
```

2.12.2 Subset Example:

```
nums <- c(10, 20, 30, 40, 50)
# Access first element: 10
nums[1]
```

```
[1] 10
```

```
# Access elements 3 to 5: 30, 40, 50
nums[3:5]
```

```
[1] 30 40 50
```

```
# Access elements 1 and 4: 10, 40
nums[c(1, 4)]
```

```
[1] 10 40
```

```
vec7 <- c(11, 23, 55, 99, 100, 500, 21, 26)
# Number of elements: 8
length(vec7)
```

```
[1] 8
```

```
# Access the last element: 26
vec7[length(vec7)]
```

```
[1] 26
```

```
# View without the first element: 23, 55, 99, 100, 500, 21, 26
vec7[-1]
```

```
[1] 23 55 99 100 500 21 26
```

```
# Subset from 3rd to 7th: 55, 99, 100, 500, 21
vec7[3:7]
```

```
[1] 55 99 100 500 21
```

```
# Up to the second last element: 55, 99, 100, 500, 21
vec7[3:(length(vec7) - 1)]
```

```
[1] 55 99 100 500 21
```

```
# Remove first three elements: 99, 100, 500, 21, 26
vec7[-(1:3)]
```

```
[1] 99 100 500 21 26
```

2.13 Vector Operations

2.13.1 Element-wise Operations:

```
vec8 <- c(12, 34, 56, 77, 78, 86, 223, 100, 45, 10)
vec9 <- c(54, 32, 87, 21, 99)
# Element-wise addition (recycling rule applies)
vec10 <- vec8 + vec9
# Element-wise division
vec11 <- vec8 / vec9
```

Explanation:

- Element-wise operations are performed with recycling if vectors are of different lengths.

2.13.2 Recycling Rule Example:

```
vec1 <- c(1, 2, 3)
vec2 <- c(4, 5)

vec_sum <- vec1 + vec2
```

Warning in vec1 + vec2: longer object length is not a multiple of shorter object length

```
# vec2 is recycled to match length of vec1: [4, 5, 4]
# Result: [5, 7, 7]
```

```
vec1 <- c(1, 2, 3)
vec2 <- c(4, 5, 6)

# Element-wise addition: [5, 7, 9]
vec_sum <- vec1 + vec2
# Scalar multiplication: [2, 4, 6]
vec_mul <- vec1 * 2
# Logical comparison: [FALSE, FALSE, TRUE]
vec_logical <- vec1 > 2
```

Explanation:

- vec_sum performs element-wise addition.
- vec_mul multiplies each element by 2.
- vec_logical creates a logical vector based on comparison.

```
vec <- c(1, 2, 3)
# Length of vector: 3
length(vec)
```

```
[1] 3
```

```
# Check if numeric: TRUE
is.numeric(vec)
```

```
[1] TRUE
```

```
# Check if character: FALSE
is.character(vec)
```

```
[1] FALSE
# Check if double: TRUE
is.double(vec)

[1] TRUE
# Check if integer: FALSE
is.integer(vec)

[1] FALSE
# Type of vector: "double"
typeof(vec)

[1] "double"
```

2.13.3 Summary

- `c()`: Combines values into a vector.
- `seq()`: Generates sequences of numbers.
- `rep()`: Repeats elements of vectors.
- `is.vector()`: Checks if an object is a vector.
- `typeof()`: Determines the type of an object.

2.14 Application Level

2.14.1 CO2 Data Analysis

```
co2 <- c(369.55, 371.14, 373.28, 375.80, 377.52, 379.80,
        381.90, 383.79, 385.60, 387.43, 389.90, 391.65,
        393.85, 396.52, 398.65, 400.83, 404.21)
year <- c(2000:2016)


# Compute the mean of CO2 values
mean(co2)

[1] 385.9659

# Compute the standard deviation of CO2 values
sd(co2)

[1] 10.6726

# Plot CO2 values against years
plot(year, co2)
```

2.14.2 Basketball Players Data Analysis

The data is form based on the data available at <https://data.world/datadavis/nba-salaries>

Instructions for this dataset:

Once executed the commands the following objects will be created:

Matrices:

1. FieldGoalAttempts
2. FieldGoals
3. Games
4. MinutesPlayed
5. Salary
6. Points
7. Players
8. Seasons

```
# Comments:  
# Seasons are labeled based on the first year in the season  
# E.g. the 2012-2013 season is preseneted as simply 2012
```

```

#Seasons
Seasons <- c("2005", "2006", "2007", "2008", "2009", "2010", "2011", "2012", "2013", "2014")

# Players
Players <- c("KobeBryant", "JoeJohnson", "LeBronJames", "CarmeloAnthony",
             "DwightHoward", "ChrisBosh", "ChrisPaul", "KevinDurant",
             "DerrickRose", "DwayneWade")

# 1. Salaries
KobeBryant_Salary <- c(15946875, 17718750, 19490625, 21262500,
                      23034375, 24806250, 25244493, 27849149, 30453805, 23500000)
JoeJohnson_Salary <- c(12000000, 12744189, 13488377, 14232567,
                       14976754, 16324500, 18038573, 19752645, 21466718, 23180790)
LeBronJames_Salary <- c(4621800, 5828090, 13041250, 14410581, 15779912,
                       1450000, 16022500, 17545000, 19067500, 20644400)
CarmeloAnthony_Salary <- c(3713640, 4694041, 13041250, 14410581,
                          15779912, 17149243, 18518574, 19450000, 22407474, 22458000)
DwightHoward_Salary <- c(4493160, 4806720, 6061274, 13758000,
                        15202590, 16647180, 18091770, 19536360, 20513178, 21436271)
ChrisBosh_Salary <- c(3348000, 4235220, 12455000, 14410581, 15779912,
                     14500000, 16022500, 17545000, 19067500, 20644400)
ChrisPaul_Salary <- c(3144240, 3380160, 3615960, 4574189, 13520500,
                     14940153, 16359805, 17779458, 18668431, 20068563)
KevinDurant_Salary <- c(0, 0, 4171200, 4484040, 4796880, 6053663,
                       15506632, 16669630, 17832627, 18995624)
DerrickRose_Salary <- c(0, 0, 0, 4822800, 5184480, 5546160,
                       6993708, 16402500, 17632688, 18862875)
DwayneWade_Salary <- c(3031920, 3841443, 13041250, 14410581, 15779912,
                      14200000, 15691000, 17182000, 18673000, 15000000)

# Matrix-1
# Step 1: Create the matrix Salary using rbind()
Salary <- rbind(KobeBryant_Salary, JoeJohnson_Salary, LeBronJames_Salary,
                CarmeloAnthony_Salary, DwightHoward_Salary, ChrisBosh_Salary,
                ChrisPaul_Salary, KevinDurant_Salary, DerrickRose_Salary, DwayneWade_Salary)

# Step 2: Remove individual player salary vectors from memory using rm()
# Purpose of rm(): This command removes the individual salary vectors from the environment,
# freeing up memory, since they have already been combined into the matrix Salary.
rm(KobeBryant_Salary, JoeJohnson_Salary, CarmeloAnthony_Salary,
    DwightHoward_Salary, ChrisBosh_Salary, LeBronJames_Salary,
    ChrisPaul_Salary, DerrickRose_Salary, DwayneWade_Salary, KevinDurant_Salary)

# Step 3: Assign column and row names to the matrix
colnames(Salary) <- Seasons
rownames(Salary) <- Players
print(Salary)

```

	2005	2006	2007	2008	2009	2010	2011
KobeBryant	15946875	17718750	19490625	21262500	23034375	24806250	25244493

JoeJohnson	12000000	12744189	13488377	14232567	14976754	16324500	18038573
LeBronJames	4621800	5828090	13041250	14410581	15779912	1450000	16022500
CarmeloAnthony	3713640	4694041	13041250	14410581	15779912	17149243	18518574
DwightHoward	4493160	4806720	6061274	13758000	15202590	16647180	18091770
ChrisBosh	3348000	4235220	12455000	14410581	15779912	14500000	16022500
ChrisPaul	3144240	3380160	3615960	4574189	13520500	14940153	16359805
KevinDurant	0	0	4171200	4484040	4796880	6053663	15506632
DerrickRose	0	0	0	4822800	5184480	5546160	6993708
DwayneWade	3031920	3841443	13041250	14410581	15779912	14200000	15691000
	2012	2013	2014				
KobeBryant	27849149	30453805	23500000				
JoeJohnson	19752645	21466718	23180790				
LeBronJames	17545000	19067500	20644400				
CarmeloAnthony	19450000	22407474	22458000				
DwightHoward	19536360	20513178	21436271				
ChrisBosh	17545000	19067500	20644400				
ChrisPaul	17779458	18668431	20068563				
KevinDurant	16669630	17832627	18995624				
DerrickRose	16402500	17632688	18862875				
DwayneWade	17182000	18673000	15000000				

2. Games

```

KobeBryant_G <- c(80,77,82,82,73,82,58,78,6,35)
JoeJohnson_G <- c(82,57,82,79,76,72,60,72,79,80)
LeBronJames_G <- c(79,78,75,81,76,79,62,76,77,69)
CarmeloAnthony_G <- c(80,65,77,66,69,77,55,67,77,40)
DwightHoward_G <- c(82,82,82,79,82,78,54,76,71,41)
ChrisBosh_G <- c(70,69,67,77,70,77,57,74,79,44)
ChrisPaul_G <- c(78,64,80,78,45,80,60,70,62,82)
KevinDurant_G <- c(35,35,80,74,82,78,66,81,81,27)
DerrickRose_G <- c(40,40,40,81,78,81,39,0,10,51)
DwayneWade_G <- c(75,51,51,79,77,76,49,69,54,62)

```

Matrix-2

```

Games <- rbind(KobeBryant_G, JoeJohnson_G, LeBronJames_G,
               CarmeloAnthony_G, DwightHoward_G, ChrisBosh_G,
               ChrisPaul_G, KevinDurant_G, DerrickRose_G, DwayneWade_G)

```

```

rm(KobeBryant_G, JoeJohnson_G, CarmeloAnthony_G,
   DwightHoward_G, ChrisBosh_G, LeBronJames_G, ChrisPaul_G,
   DerrickRose_G, DwayneWade_G, KevinDurant_G)

```

```
colnames(Games) <- Seasons
```

3. Minutes Played

```

KobeBryant_MP <- c(3277,3140,3192,2960,2835,2779,2232,3013,177,1207)
JoeJohnson_MP <- c(3340,2359,3343,3124,2886,2554,2127,2642,2575,2791)
LeBronJames_MP <- c(3361,3190,3027,3054,2966,3063,2326,2877,2902,2493)
CarmeloAnthony_MP <- c(2941,2486,2806,2277,2634,2751,1876,2482,2982,1428)
DwightHoward_MP <- c(3021,3023,3088,2821,2843,2935,2070,2722,2396,1223)

```

```
ChrisBosh_MP <- c(2751,2658,2425,2928,2526,2795,2007,2454,2531,1556)
ChrisPaul_MP <- c(2808,2353,3006,3002,1712,2880,2181,2335,2171,2857)
KevinDurant_MP <- c(1255,1255,2768,2885,3239,3038,2546,3119,3122,913)
DerrickRose_MP <- c(1168,1168,1168,3000,2871,3026,1375,0,311,1530)
DwayneWade_MP <- c(2892,1931,1954,3048,2792,2823,1625,2391,1775,1971)
```

Matrix-3

```
MinutesPlayed <- rbind(KobeBryant_MP, JoeJohnson_MP,
                       LeBronJames_MP, CarmeloAnthony_MP,
                       DwightHoward_MP, ChrisBosh_MP, ChrisPaul_MP,
                       KevinDurant_MP, DerrickRose_MP, DwayneWade_MP)
```

```
rm(KobeBryant_MP, JoeJohnson_MP, CarmeloAnthony_MP,
   DwightHoward_MP, ChrisBosh_MP, LeBronJames_MP, ChrisPaul_MP,
   DerrickRose_MP, DwayneWade_MP, KevinDurant_MP)
```

```
colnames(MinutesPlayed) <- Seasons
```

```
rownames(MinutesPlayed) <- Players
```

4. Field Goals

```
KobeBryant_FG <- c(978,813,775,800,716,740,574,738,31,266)
JoeJohnson_FG <- c(632,536,647,620,635,514,423,445,462,446)
LeBronJames_FG <- c(875,772,794,789,768,758,621,765,767,624)
CarmeloAnthony_FG <- c(756,691,728,535,688,684,441,669,743,358)
DwightHoward_FG <- c(468,526,583,560,510,619,416,470,473,251)
ChrisBosh_FG <- c(549,543,507,615,600,524,393,485,492,343)
ChrisPaul_FG <- c(407,381,630,631,314,430,425,412,406,568)
KevinDurant_FG <- c(306,306,587,661,794,711,643,731,849,238)
DerrickRose_FG <- c(208,208,208,574,672,711,302,0,58,338)
DwayneWade_FG <- c(699,472,439,854,719,692,416,569,415,509)
```

Matrix-4

```
FieldGoals <- rbind(KobeBryant_FG, JoeJohnson_FG, LeBronJames_FG,
                    CarmeloAnthony_FG, DwightHoward_FG, ChrisBosh_FG,
                    ChrisPaul_FG, KevinDurant_FG, DerrickRose_FG, DwayneWade_FG)
```

```
rm(KobeBryant_FG, JoeJohnson_FG, LeBronJames_FG,
   CarmeloAnthony_FG, DwightHoward_FG, ChrisBosh_FG,
   ChrisPaul_FG, KevinDurant_FG, DerrickRose_FG, DwayneWade_FG)
```

```
colnames(FieldGoals) <- Seasons
```

```
rownames(FieldGoals) <- Players
```

5. Field Goal Attempts

```
KobeBryant_FGA <- c(2173,1757,1690,1712,1569,1639,1336,1595,73,713)
JoeJohnson_FGA <- c(1395,1139,1497,1420,1386,1161,931,1052,1018,1025)
LeBronJames_FGA <- c(1823,1621,1642,1613,1528,1485,1169,1354,1353,1279)
```

```

CarmeloAnthony_FGA <- c(1572,1453,1481,1207,1502,1503,1025,1489,1643,806)
DwightHoward_FGA <- c(881,873,974,979,834,1044,726,813,800,423)
ChrisBosh_FGA <- c(1087,1094,1027,1263,1158,1056,807,907,953,745)
ChrisPaul_FGA <- c(947,871,1291,1255,637,928,890,856,870,1170)
KevinDurant_FGA <- c(647,647,1366,1390,1668,1538,1297,1433,1688,467)
DerrickRose_FGA <- c(436,436,436,1208,1373,1597,695,0,164,835)
DwayneWade_FGA <- c(1413,962,937,1739,1511,1384,837,1093,761,1084)

# Matrix-5
FieldGoalAttempts <- rbind(KobeBryant_FGA, JoeJohnson_FGA,
                           LeBronJames_FGA, CarmeloAnthony_FGA,
                           DwightHoward_FGA, ChrisBosh_FGA, ChrisPaul_FGA,
                           KevinDurant_FGA, DerrickRose_FGA, DwayneWade_FGA)

rm(KobeBryant_FGA, JoeJohnson_FGA, LeBronJames_FGA,
   CarmeloAnthony_FGA, DwightHoward_FGA, ChrisBosh_FGA,
   ChrisPaul_FGA, KevinDurant_FGA, DerrickRose_FGA, DwayneWade_FGA)

colnames(FieldGoalAttempts) <- Seasons
rownames(FieldGoalAttempts) <- Players

# 6.Points
KobeBryant_PTS <- c(2832,2430,2323,2201,1970,2078,1616,2133,83,782)
JoeJohnson_PTS <- c(1653,1426,1779,1688,1619,1312,1129,1170,1245,1154)
LeBronJames_PTS <- c(2478,2132,2250,2304,2258,2111,1683,2036,2089,1743)
CarmeloAnthony_PTS <- c(2122,1881,1978,1504,1943,1970,1245,1920,2112,966)
DwightHoward_PTS <- c(1292,1443,1695,1624,1503,1784,1113,1296,1297,646)
ChrisBosh_PTS <- c(1572,1561,1496,1746,1678,1438,1025,1232,1281,928)
ChrisPaul_PTS <- c(1258,1104,1684,1781,841,1268,1189,1186,1185,1564)
KevinDurant_PTS <- c(903,903,1624,1871,2472,2161,1850,2280,2593,686)
DerrickRose_PTS <- c(597,597,597,1361,1619,2026,852,0,159,904)
DwayneWade_PTS <- c(2040,1397,1254,2386,2045,1941,1082,1463,1028,1331)

# Matrix-6
Points <- rbind(KobeBryant_PTS, JoeJohnson_PTS, LeBronJames_PTS,
                CarmeloAnthony_PTS, DwightHoward_PTS, ChrisBosh_PTS,
                ChrisPaul_PTS, KevinDurant_PTS, DerrickRose_PTS, DwayneWade_PTS)

rm(KobeBryant_PTS, JoeJohnson_PTS, LeBronJames_PTS,
   CarmeloAnthony_PTS, DwightHoward_PTS, ChrisBosh_PTS,
   ChrisPaul_PTS, KevinDurant_PTS, DerrickRose_PTS, DwayneWade_PTS)

colnames(Points) <- Seasons
rownames(Points) <- Players

```

2.15 Questions

2.15.1 How many games did ChrisPaul play in 2011?

```
no_of_games_CP = Games["ChrisPaul_G", "2011"]
paste(no_of_games_CP, "no. of games ChrisPaul play in 2011")
```

```
[1] "60 no. of games ChrisPaul play in 2011"
```

Explanation:

- Games is a matrix where the rows are players and the columns are seasons. To get the number of games Chris Paul played in 2011, you access the element in the row labeled “ChrisPaul” and the column labeled “2011”.
- paste() combines the result with a string to provide a readable output.

2.15.2 What are the field goals per game for each player?

```
dim(Games)
```

```
[1] 10 10
```

```
dim(FieldGoals) # check dimension
```

```
[1] 10 10
```

```
FieldGoals_Per_Game = FieldGoals / Games
print(FieldGoals_Per_Game)
```

	2005	2006	2007	2008	2009	2010
KobeBryant	12.225000	10.558442	9.451220	9.756098	9.808219	9.024390
JoeJohnson	7.707317	9.403509	7.890244	7.848101	8.355263	7.138889
LeBronJames	11.075949	9.897436	10.586667	9.740741	10.105263	9.594937
CarmeloAnthony	9.450000	10.630769	9.454545	8.106061	9.971014	8.883117
DwightHoward	5.707317	6.414634	7.109756	7.088608	6.219512	7.935897
ChrisBosh	7.842857	7.869565	7.567164	7.987013	8.571429	6.805195
ChrisPaul	5.217949	5.953125	7.875000	8.089744	6.977778	5.375000
KevinDurant	8.742857	8.742857	7.337500	8.932432	9.682927	9.115385
DerrickRose	5.200000	5.200000	5.200000	7.086420	8.615385	8.777778
DwayneWade	9.320000	9.254902	8.607843	10.810127	9.337662	9.105263
	2011	2012	2013	2014		
KobeBryant	9.896552	9.461538	5.166667	7.600000		
JoeJohnson	7.050000	6.180556	5.848101	5.575000		
LeBronJames	10.016129	10.065789	9.961039	9.043478		
CarmeloAnthony	8.018182	9.985075	9.649351	8.950000		
DwightHoward	7.703704	6.184211	6.661972	6.121951		
ChrisBosh	6.894737	6.554054	6.227848	7.795455		
ChrisPaul	7.083333	5.885714	6.548387	6.926829		
KevinDurant	9.742424	9.024691	10.481481	8.814815		
DerrickRose	7.743590	NaN	5.800000	6.627451		
DwayneWade	8.489796	8.246377	7.685185	8.209677		

```
round(FieldGoals_Per_Game)
```

	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014
KobeBryant	12	11	9	10	10	9	10	9	5	8
JoeJohnson	8	9	8	8	8	7	7	6	6	6
LeBronJames	11	10	11	10	10	10	10	10	10	9
CarmeloAnthony	9	11	9	8	10	9	8	10	10	9
DwightHoward	6	6	7	7	6	8	8	6	7	6
ChrisBosh	8	8	8	8	9	7	7	7	6	8

ChrisPaul	5	6	8	8	7	5	7	6	7	7
KevinDurant	9	9	7	9	10	9	10	9	10	9
DerrickRose	5	5	5	7	9	9	8	NaN	6	7
DwayneWade	9	9	9	11	9	9	8	8	8	8

```
round(FieldGoals_Per_Game, 1) # with one decimal
```

	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014
KobeBryant	12.2	10.6	9.5	9.8	9.8	9.0	9.9	9.5	5.2	7.6
JoeJohnson	7.7	9.4	7.9	7.8	8.4	7.1	7.0	6.2	5.8	5.6
LeBronJames	11.1	9.9	10.6	9.7	10.1	9.6	10.0	10.1	10.0	9.0
CarmeloAnthony	9.4	10.6	9.5	8.1	10.0	8.9	8.0	10.0	9.6	8.9
DwightHoward	5.7	6.4	7.1	7.1	6.2	7.9	7.7	6.2	6.7	6.1
ChrisBosh	7.8	7.9	7.6	8.0	8.6	6.8	6.9	6.6	6.2	7.8
ChrisPaul	5.2	6.0	7.9	8.1	7.0	5.4	7.1	5.9	6.5	6.9
KevinDurant	8.7	8.7	7.3	8.9	9.7	9.1	9.7	9.0	10.5	8.8
DerrickRose	5.2	5.2	5.2	7.1	8.6	8.8	7.7	NaN	5.8	6.6
DwayneWade	9.3	9.3	8.6	10.8	9.3	9.1	8.5	8.2	7.7	8.2

Explanation:

- FieldGoals and Games are matrices of the same dimensions.
- The division FieldGoals / Games computes the field goals per game for each player by dividing the number of field goals by the number of games played.
- round() is used to round the results to a specified number of decimal places for better readability.

2.15.3 How many minutes did each player play per game?

```
Min_ply_per_game = MinutesPlayed / Games
print(Min_ply_per_game)
```

	2005	2006	2007	2008	2009	2010	2011
KobeBryant	40.96250	40.77922	38.92683	36.09756	38.83562	33.89024	38.48276
JoeJohnson	40.73171	41.38596	40.76829	39.54430	37.97368	35.47222	35.45000
LeBronJames	42.54430	40.89744	40.36000	37.70370	39.02632	38.77215	37.51613
CarmeloAnthony	36.76250	38.24615	36.44156	34.50000	38.17391	35.72727	34.10909
DwightHoward	36.84146	36.86585	37.65854	35.70886	34.67073	37.62821	38.33333
ChrisBosh	39.30000	38.52174	36.19403	38.02597	36.08571	36.29870	35.21053
ChrisPaul	36.00000	36.76562	37.57500	38.48718	38.04444	36.00000	36.35000
KevinDurant	35.85714	35.85714	34.60000	38.98649	39.50000	38.94872	38.57576
DerrickRose	29.20000	29.20000	29.20000	37.03704	36.80769	37.35802	35.25641
DwayneWade	38.56000	37.86275	38.31373	38.58228	36.25974	37.14474	33.16327

	2012	2013	2014
KobeBryant	38.62821	29.50000	34.48571
JoeJohnson	36.69444	32.59494	34.88750
LeBronJames	37.85526	37.68831	36.13043
CarmeloAnthony	37.04478	38.72727	35.70000
DwightHoward	35.81579	33.74648	29.82927
ChrisBosh	33.16216	32.03797	35.36364
ChrisPaul	33.35714	35.01613	34.84146
KevinDurant	38.50617	38.54321	33.81481
DerrickRose	NaN	31.10000	30.00000
DwayneWade	34.65217	32.87037	31.79032

Explanation:

- MinutesPlayed and Games are matrices with the same dimensions. Dividing MinutesPlayed by Games gives the average number of minutes played per game for each player.

2.15.4 How much is per minute worth for each player?

```
Per_min_worth = Salary / MinutesPlayed
print(Per_min_worth)
```

	2005	2006	2007	2008	2009	2010	2011
KobeBryant	4866.303	5642.914	6106.086	7183.277	8125.000	8926.3224	11310.257
JoeJohnson	3592.814	5402.369	4034.812	4555.879	5189.450	6391.7384	8480.758
LeBronJames	1375.126	1826.987	4308.309	4718.592	5320.267	473.3921	6888.435
CarmeloAnthony	1262.713	1888.190	4647.630	6328.758	5990.855	6233.8215	9871.308
DwightHoward	1487.309	1590.050	1962.848	4876.994	5347.376	5671.9523	8739.986
ChrisBosh	1217.012	1593.386	5136.082	4921.647	6246.996	5187.8354	7983.308
ChrisPaul	1119.744	1436.532	1202.914	1523.714	7897.488	5187.5531	7501.057
KevinDurant	0.000	0.000	1506.936	1554.260	1480.976	1992.6475	6090.586
DerrickRose	0.000	0.000	0.000	1607.600	1805.810	1832.8354	5086.333
DwayneWade	1048.382	1989.354	6674.130	4727.881	5651.831	5030.1098	9656.000
	2012	2013	2014				
KobeBryant	9242.997	172055.395	19469.760				
JoeJohnson	7476.399	8336.590	8305.550				
LeBronJames	6098.366	6570.469	8280.947				
CarmeloAnthony	7836.422	7514.243	15726.891				
DwightHoward	7177.208	8561.427	17527.613				
ChrisBosh	7149.552	7533.584	13267.609				
ChrisPaul	7614.329	8599.001	7024.348				
KevinDurant	5344.543	5711.924	20805.722				
DerrickRose	Inf	56696.746	12328.676				
DwayneWade	7186.115	10520.000	7610.350				

Explanation:

- Salary and MinutesPlayed are matrices. Dividing Salary by MinutesPlayed computes the worth per minute for each player, which tells you how much salary is earned per minute of play.

2.15.5 How accurate is each player?

```
Accurate <- FieldGoals / FieldGoalAttempts
print(Accurate)
```

	2005	2006	2007	2008	2009	2010
KobeBryant	0.4500690	0.4627205	0.4585799	0.4672897	0.4563416	0.4514948
JoeJohnson	0.4530466	0.4705882	0.4321977	0.4366197	0.4581530	0.4427218
LeBronJames	0.4799781	0.4762492	0.4835566	0.4891507	0.5026178	0.5104377
CarmeloAnthony	0.4809160	0.4755678	0.4915598	0.4432477	0.4580559	0.4550898
DwightHoward	0.5312145	0.6025200	0.5985626	0.5720123	0.6115108	0.5929119
ChrisBosh	0.5050598	0.4963437	0.4936709	0.4869359	0.5181347	0.4962121
ChrisPaul	0.4297782	0.4374282	0.4879938	0.5027888	0.4929356	0.4633621
KevinDurant	0.4729521	0.4729521	0.4297218	0.4755396	0.4760192	0.4622887
DerrickRose	0.4770642	0.4770642	0.4770642	0.4751656	0.4894392	0.4452098
DwayneWade	0.4946921	0.4906445	0.4685165	0.4910868	0.4758438	0.5000000
	2011	2012	2013	2014		
KobeBryant	0.4296407	0.4626959	0.4246575	0.3730715		
JoeJohnson	0.4543502	0.4230038	0.4538310	0.4351220		

LeBronJames	0.5312233	0.5649926	0.5668884	0.4878812
CarmeloAnthony	0.4302439	0.4492948	0.4522215	0.4441687
DwightHoward	0.5730028	0.5781058	0.5912500	0.5933806
ChrisBosh	0.4869888	0.5347299	0.5162644	0.4604027
ChrisPaul	0.4775281	0.4813084	0.4666667	0.4854701
KevinDurant	0.4957594	0.5101186	0.5029621	0.5096360
DerrickRose	0.4345324	NaN	0.3536585	0.4047904
DwayneWade	0.4970131	0.5205855	0.5453351	0.4695572

Explanation:

- FieldGoals and FieldGoalAttempts are matrices. Dividing FieldGoals by FieldGoalAttempts calculates the shooting accuracy (or field goal percentage) for each player, which measures how often a player makes a field goal attempt.

2.15.6 Who is good at 3-pointers?

```
# Calculate total points for each player
r_sum <- matrix(rowSums(Points), nrow = length(Players), dimnames = list(Players, "Total Points"))
print(r_sum)
```

	Total Points
KobeBryant	18448
JoeJohnson	14175
LeBronJames	21084
CarmeloAnthony	17641
DwightHoward	13693
ChrisBosh	13957
ChrisPaul	13060
KevinDurant	17343
DerrickRose	8712
DwayneWade	15967

```
# Sort players by total points and print the sorted results
sorted_r_sum <- sort(r_sum[, "Total Points"], decreasing = TRUE)
print(sorted_r_sum)
```

LeBronJames	KobeBryant	CarmeloAnthony	KevinDurant	DwayneWade
21084	18448	17641	17343	15967
JoeJohnson	ChrisBosh	DwightHoward	ChrisPaul	DerrickRose
14175	13957	13693	13060	8712

```
# Print the highest ranked player
highest_ranked_player <- names(sorted_r_sum)[1]
highest_ranked_points <- sorted_r_sum[1]
cat("The highest ranked player based on total points is:",
    highest_ranked_player, "with",
    highest_ranked_points, "total points.\n")
```

The highest ranked player based on total points is: LeBronJames with 21084 total points.

- rowSums(Points): This calculates the sum of points scored by each player across all seasons.
- matrix(..., nrow = length(Players), dimnames = list(Players, "Total Points")): This converts the total points into a matrix. Each row corresponds to a player, and the column is named "Total Points".
- print(r_sum): This prints the matrix r_sum to the console, showing each player and their total points.

- `names(sorted_r_sum)[1]`: This extracts the name of the player with the highest total points (the first element in `sorted_r_sum`).
- `sorted_r_sum[1]`: This extracts the highest total points value.
- `cat(...)`: This concatenates and prints a message to the console stating who the highest-ranked player is and their total points.

3 Part-3

3.1 Read the Data from a CSV File

```
mydata <- read.csv("datasets/DemographicsData.csv")
```

- `read.csv("datasets/DemographicsData.csv")`: Reads the CSV file into a data frame named `mydata`.

`mydata`

- `mydata`: Displays the entire data frame.

3.2 Explore and Understand the Data

```
# Number of rows in the data frame  
nrow(mydata)
```

```
[1] 195
```

```
# Number of columns in the data frame  
ncol(mydata)
```

```
[1] 5
```

```
# Dimensions of the data frame (rows, columns)  
dim(mydata)
```

```
[1] 195 5
```

```
# Number of rows  
dim(mydata)[1]
```

```
[1] 195
```

- `nrow(mydata)`: Returns the number of rows.
- `ncol(mydata)`: Returns the number of columns.
- `dim(mydata)`: Returns a vector with the number of rows and columns.
- `dim(mydata)[1]`: Extracts the number of rows from the dimensions vector.

```
# View the first 6 rows of the data frame  
head(mydata)
```

	Country.Name	Country.Code	Birth.rate	Internet.users
1	Aruba	ABW	10.244	78.9
2	Afghanistan	AFG	35.253	5.9
3	Angola	AGO	45.985	19.1
4	Albania	ALB	12.877	57.2
5	United Arab Emirates	ARE	11.044	88.0
6	Argentina	ARG	17.716	59.9
	Income.Group			
1	High income			
2	Low income			
3	Upper middle income			
4	Upper middle income			
5	High income			
6	High income			

```
# View the last 6 rows of the data frame
tail(mydata)
```

	Country.Name	Country.Code	Birth.rate	Internet.users	Income.Group
190	Samoa	WSM	26.172	15.3	Lower middle income
191	Yemen, Rep.	YEM	32.947	20.0	Lower middle income
192	South Africa	ZAF	20.850	46.5	Upper middle income
193	Congo, Dem. Rep.	COD	42.394	2.2	Low income
194	Zambia	ZMB	40.471	15.4	Lower middle income
195	Zimbabwe	ZWE	35.715	18.5	Low income

- **head(mydata)**: Shows the first 6 rows.
- **tail(mydata)**: Shows the last 6 rows.

```
# Structure of the data frame, showing column types and a preview
str(mydata)
```

```
'data.frame': 195 obs. of 5 variables:
 $ Country.Name : chr "Aruba" "Afghanistan" "Angola" "Albania" ...
 $ Country.Code : chr "ABW" "AFG" "AGO" "ALB" ...
 $ Birth.rate : num 10.2 35.3 46 12.9 11 ...
 $ Internet.users: num 78.9 5.9 19.1 57.2 88 ...
 $ Income.Group : chr "High income" "Low income" "Upper middle income" "Upper middle income" ...
```

```
# Summary statistics for each column
summary(mydata)
```

Country.Name	Country.Code	Birth.rate	Internet.users
Length:195	Length:195	Min. : 7.90	Min. : 0.90
Class :character	Class :character	1st Qu.:12.12	1st Qu.:14.52
Mode :character	Mode :character	Median :19.68	Median :41.00
		Mean :21.47	Mean :42.08
		3rd Qu.:29.76	3rd Qu.:66.22
		Max. :49.66	Max. :96.55

Income.Group
Length:195
Class :character
Mode :character

- **str(mydata)**: Provides the structure of the data frame, including data types and a preview of the data.
- **summary(mydata)**: Provides summary statistics for each column, such as min, max, mean, median, etc.

3.3 Subsetting the Data Frame

```
# Extract the first row as a data frame
mydata[1, ]
```

	Country.Name	Country.Code	Birth.rate	Internet.users	Income.Group
1	Aruba	ABW	10.244	78.9	High income

```
# Check if the result is still a data frame
is.data.frame(mydata[1, ])
```

```
[1] TRUE
```

- `mydata[1,]`: Extracts the first row of `mydata` as a data frame.
- `is.data.frame(mydata[1,])`: Checks if the extracted row is a data frame (it is).

```
# Extract the first column as a vector
mydata[, 1]
```

```
# Check if the result is a data frame (it's not)
is.data.frame(mydata[, 1])
```

```
[1] FALSE
```

```
# Confirm the result is a vector
is.vector(mydata[, 1])
```

```
[1] TRUE
```

```
# Extract the first column as a data frame (TRUE)
is.data.frame(mydata[, 1, drop = F])
```

```
[1] TRUE
```

- `mydata[, 1]`: Extracts the first column as a vector.
- `is.data.frame(mydata[, 1])`: Checks if the extracted column is a data frame (it's not; it's a vector).
- `is.vector(mydata[, 1])`: Confirms that the result is a vector.
- `mydata[, 1, drop = F]`: Extracts the first column as a data frame (using `drop = FALSE` to preserve the data frame structure).
- **Without `drop = FALSE`**: If you simply use `mydata[, 1]`, you get a vector of the first column if `mydata` has only one column selected.
- **With `drop = FALSE`**: Using `mydata[, 1, drop = FALSE]` ensures that the result is a data frame with one column, even if only a single column is extracted.

3.4 Accessing Columns Using `$`

```
# Get 'Country.Name' column
country_name = mydata$Country.Name
head(country_name, 10)
```

```
[1] "Aruba"           "Afghanistan"      "Angola"
[4] "Albania"         "United Arab Emirates" "Argentina"
[7] "Armenia"         "Antigua and Barbuda" "Australia"
[10] "Austria"
```

```
# First 10 rows
mydata[1:10,]
```

	Country.Name	Country.Code	Birth.rate	Internet.users
1	Aruba	ABW	10.244	78.9000
2	Afghanistan	AFG	35.253	5.9000
3	Angola	AGO	45.985	19.1000
4	Albania	ALB	12.877	57.2000
5	United Arab Emirates	ARE	11.044	88.0000
6	Argentina	ARG	17.716	59.9000
7	Armenia	ARM	13.308	41.9000

```

8   Antigua and Barbuda      ATG      16.447      63.4000
9           Australia        AUS      13.200      83.0000
10          Austria          AUT       9.400      80.6188
    Income.Group
1           High income
2           Low income
3   Upper middle income
4   Upper middle income
5           High income
6           High income
7   Lower middle income
8           High income
9           High income
10          High income

```

```

# Extract the 5th and 99th rows from the dataframe
mydata[c(5, 99), ]

```

```

      Country.Name Country.Code Birth.rate Internet.users
5   United Arab Emirates      ARE      11.044          88.0
99          Lebanon          LBN      13.426          70.5
    Income.Group
5           High income
99 Upper middle income

```

```

# Extract the 3rd and 55th rows from the dataframe
mydata[c(3, 55), ]

```

```

      Country.Name Country.Code Birth.rate Internet.users      Income.Group
3           Angola          AGO      45.985          19.1 Upper middle income
55          Estonia          EST      10.300          79.4          High income

```

Explanation:

- `mydata$Country.Name`: Extracts the 'Country.Name' column as a vector.
- `mydata[1:10,]`: Displays the first 10 rows of the dataframe.
- `mydata[c(5, 99),]`: Shows rows 5 and 99.
- `mydata[c(3, 55),]`: Shows rows 3 and 55.

3.5 Create an Additional Attribute (Add a Column)

```

# Add 'dummy' column as product of 'Birth.rate' and 'Internet.users'
mydata$dummy <- mydata$Birth.rate * mydata$Internet.users
head(mydata, 5)

```

```

      Country.Name Country.Code Birth.rate Internet.users
1           Aruba          ABW      10.244          78.9
2   Afghanistan          AFG      35.253           5.9
3           Angola          AGO      45.985          19.1
4           Albania          ALB      12.877          57.2
5   United Arab Emirates      ARE      11.044          88.0
    Income.Group      dummy
1   High income 808.2516
2   Low income 207.9927
3   Upper middle income 878.3135

```

```
4 Upper middle income 736.5644
5 High income 971.8720
```

```
# Remove 'dummy' column
mydata$dummy <- NULL
```

Explanation:

- `mydata$dummy <- mydata$Birth.rate * mydata$Internet.users`: Adds a new column 'dummy' which is the product of 'Birth.rate' and 'Internet.users'.
- `mydata$dummy <- NULL`: Removes the 'dummy' column from the dataframe.

3.6 Filter the Data

```
# Filter rows where 'Internet.users' < 2
myfilter1 <- mydata$Internet.users < 2
```

```
# Show filtered rows
mydata[myfilter1, ]
```

	Country.Name	Country.Code	Birth.rate	Internet.users	Income.Group
12	Burundi	BDI	44.151	1.3	Low income
53	Eritrea	ERI	34.800	0.9	Low income
56	Ethiopia	ETH	32.925	1.9	Low income
65	Guinea	GIN	37.337	1.6	Low income
118	Myanmar	MMR	18.119	1.6	Lower middle income
128	Niger	NER	49.661	1.7	Low income
155	Sierra Leone	SLE	36.729	1.7	Low income
157	Somalia	SOM	43.891	1.5	Low income
173	Timor-Leste	TLS	35.755	1.1	Lower middle income

```
# Count filtered rows
nrow(mydata[myfilter1, ])
```

```
[1] 9
```

```
# Filter rows where 'Internet.users' < 4 and 'Birth.rate' > 40
myfilter2 <- mydata$Internet.users < 4 & mydata$Birth.rate > 40
```

```
# Show filtered rows
mydata[myfilter2, ]
```

	Country.Name	Country.Code	Birth.rate	Internet.users	Income.Group
12	Burundi	BDI	44.151	1.3	Low income
116	Mali	MLI	44.138	3.5	Low income
128	Niger	NER	49.661	1.7	Low income
157	Somalia	SOM	43.891	1.5	Low income
168	Chad	TCD	45.745	2.3	Low income
193	Congo, Dem. Rep.	COD	42.394	2.2	Low income

```
# Get row for 'New Zealand'
myfilter3 <- mydata$Country.Name == "New Zealand"
```

```
# Show row for 'New Zealand'
mydata[myfilter3, ]
```

	Country.Name	Country.Code	Birth.rate	Internet.users	Income.Group
--	--------------	--------------	------------	----------------	--------------

134	New Zealand	NZL	13.12	82.78	High income
-----	-------------	-----	-------	-------	-------------

```
# Another way to get row for 'New Zealand'
mydata[mydata$Country.Name == "New Zealand", ]
```

	Country.Name	Country.Code	Birth.rate	Internet.users	Income.Group
134	New Zealand	NZL	13.12	82.78	High income

Explanation:

- `mydata$Internet.users < 2`: Creates a logical vector for rows where 'Internet.users' is less than 2.
- `mydata[myfilter1,]`: Shows rows where the condition is true.
- `nrow(mydata[myfilter1,])`: Counts how many rows meet the condition.
- `mydata$Internet.users < 4 & mydata$Birth.rate > 40`: Creates a logical vector for rows where 'Internet.users' is less than 4 and 'Birth.rate' is greater than 40.
- `mydata[myfilter2,]`: Shows rows where both conditions are true.
- `mydata$Country.Name == "New Zealand"`: Creates a logical vector to filter rows where 'Country.Name' is "New Zealand".
- `mydata[myfilter3,]`: Shows the row for "New Zealand".
- `mydata[mydata$Country.Name == "New Zealand",]`: Another way to filter rows for "New Zealand".

3.7 Creating a New Attribute Based on Conditions

```
# Add 'InternetLevel' column with default "Low"
mydata$InternetLevel <- "Low"
```

```
# Set "High" for 'Internet.users' >= 70
mydata[mydata$Internet.users >= 70, "InternetLevel"] <- "High"
```

```
# Set "Medium" for 'Internet.users' between 40 and 69
mydata[mydata$Internet.users < 70 & mydata$Internet.users >= 40, "InternetLevel"] <- "Medium"
```

```
# Set "Low" for 'Internet.users' < 40
mydata[mydata$Internet.users < 40, "InternetLevel"] <- "Low"
```

```
head(mydata,5)
```

	Country.Name	Country.Code	Birth.rate	Internet.users
1	Aruba	ABW	10.244	78.9
2	Afghanistan	AFG	35.253	5.9
3	Angola	AGO	45.985	19.1
4	Albania	ALB	12.877	57.2
5	United Arab Emirates	ARE	11.044	88.0

	Income.Group	InternetLevel
1	High income	High
2	Low income	Low
3	Upper middle income	Low
4	Upper middle income	Medium
5	High income	High

Explanation:

- `mydata$InternetLevel <- "Low"`: Adds a new column 'InternetLevel' with the default value "Low".

- `mydata[mydata$Internet.users >= 70, "InternetLevel"] <- "High"`: Sets 'InternetLevel' to "High" where 'Internet.users' is 70 or more.
- `mydata[mydata$Internet.users < 70 & mydata$Internet.users >= 40, "InternetLevel"] <- "Medium"`: Sets 'InternetLevel' to "Medium" where 'Internet.users' is between 40 and 69.
- `mydata[mydata$Internet.users < 40, "InternetLevel"] <- "Low"`: Ensures 'InternetLevel' is "Low" where 'Internet.users' is less than 40.

3.8 Quick Plotting with qplot

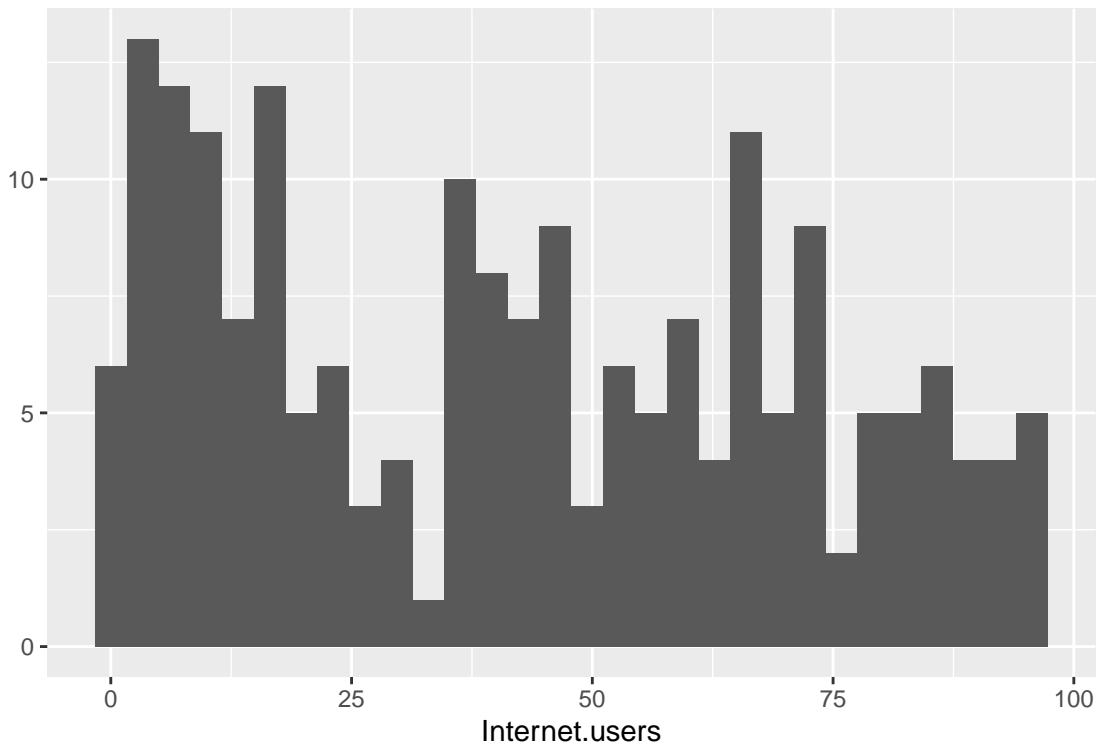
```
# Load ggplot2
library(ggplot2)

# Load data again
mydata <- read.csv("datasets/DemographicsData.csv")


# Histogram of 'Internet.users'
qplot(data = mydata, x = Internet.users)
```

Warning: `qplot()` was deprecated in ggplot2 3.4.0.
 This warning is displayed once every 8 hours.
 Call ``lifecycle::last_lifecycle_warnings()`` to see where this warning was generated.


``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.




```
# Scatter plot of 'Internet.users' vs 'Birth.rate'
qplot(data = mydata, x = Internet.users, y = Birth.rate)
```



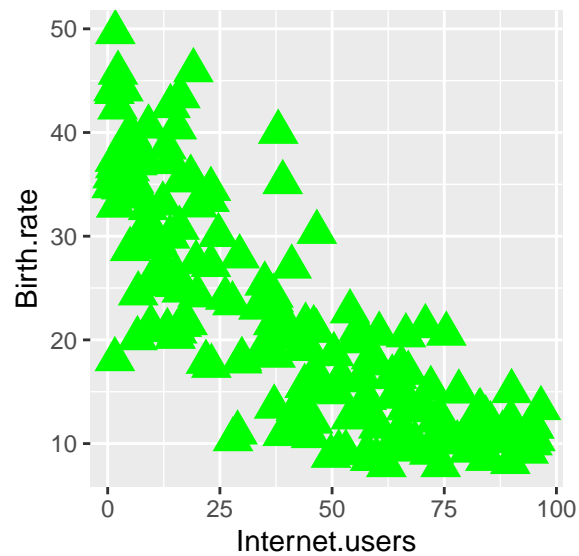
```
# Scatter plot with size
qplot(data = mydata, x = Internet.users, y = Birth.rate, size = I(5))
```




```
# Scatter plot with color
qplot(data = mydata, x = Internet.users, y = Birth.rate, size = I(5), color = I("brown"))
```




```
# Scatter plot with shape
qplot(data = mydata, x = Internet.users, y = Birth.rate,
      size = I(5), color = I("green"), pch = I(17))
```



```
# Scatter plot with transparency
qplot(data = mydata, x = Internet.users, y = Birth.rate,
      size = I(5), color = I("brown"),
      pch = I(17), alpha = I(0.5))
```



```
# Scatter plot with color by 'Income.Group'
qplot(data = mydata, x = Internet.users, y = Birth.rate,
      size = I(5), color = Income.Group,
      pch = I(19), alpha = I(0.5))
```



Explanation:

- `library(ggplot2)`: Loads the ggplot2 library for plotting.
- `?qplot()`: Shows help for the `qplot` function.
- `mydata <- read.csv("DemographicsData.csv")`: Reloads the data.

- `qplot(data = mydata, x = Internet.users)`: Creates a histogram of the 'Internet.users' variable.
- `qplot(data = mydata, x = Internet.users, y = Birth.rate)`: Creates a scatter plot of 'Internet.users' versus 'Birth.rate'.
- `qplot(data = mydata, x = Internet.users, y = Birth.rate, size = I(5))`: Adjusts the size of the plot points.
- `qplot(data = mydata, x = Internet.users, y = Birth.rate, size = I(5), color = I("brown"))`: Changes the color of the points.
- `qplot(data = mydata, x = Internet.users, y = Birth.rate, size = I(5), color = I("green"), pch = I(17))`: Changes the shape of the points.
- `qplot(data = mydata, x = Internet.users, y = Birth.rate, size = I(5), color = I("brown"), pch = I(17), alpha = I(0.5))`: Adjusts the transparency of the points.
- `qplot(data = mydata, x = Internet.users, y = Birth.rate, size = I(5), color = Income.Group, pch = I(19), alpha = I(0.5))`: Colors points based on 'Income.Group' and adjusts transparency.

4 Part-4

4.1 Reading the CSV Files

```
mydf1 <- read.csv("datasets/DemographicsData.csv")
mydf2 <- read.csv("datasets/CountryRegion.csv")
```

- `read.csv("DemographicsData.csv")`: Reads the CSV file named "DemographicsData.csv" into a dataframe `mydf1`.
- `read.csv("CountryRegion.csv")`: Reads the CSV file named "CountryRegion.csv" into a dataframe `mydf2`.

4.2 Exploring the Dataframes

```
# Check column names of the first dataframe
colnames(mydf1)
```

```
[1] "Country.Name" "Country.Code" "Birth.rate" "Internet.users"
[5] "Income.Group"
```

```
# Check column names of the second dataframe
colnames(mydf2)
```

```
[1] "Countries_2021_Dataset" "Codes_2021_Dataset" "Regions_2021_Dataset"
```

```
# Merge the two dataframes by matching 'Country.Code'
# in mydf1 with 'Codes_2021_Dataset' in mydf2
```

```
mymerg <- merge(mydf1, mydf2, by.x = "Country.Code", by.y = "Codes_2021_Dataset")
head(mymerg, 5)
```

	Country.Code	Country.Name	Birth.rate	Internet.users
1	ABW	Aruba	10.244	78.9
2	AFG	Afghanistan	35.253	5.9
3	AGO	Angola	45.985	19.1
4	ALB	Albania	12.877	57.2
5	ARE	United Arab Emirates	11.044	88.0

	Income.Group	Countries_2021_Dataset	Regions_2021_Dataset
1	High income	Aruba	The Americas
2	Low income	Afghanistan	Asia
3	Upper middle income	Angola	Africa
4	Upper middle income	Albania	Europe
5	High income	United Arab Emirates	Middle East

```
# Remove the unnecessary column 'Countries_2021_Dataset' from the merged dataframe
mymerg$Countries_2021_Dataset <- NULL
```

```
# Save the merged dataframe to a new CSV file without row names
write.csv(mymerg, "merged.csv", row.names = FALSE)
```

4.2.1 Explanation:

- `mydf1 <- read.csv("DemographicsData.csv")`: Loads the first dataset from the CSV file into the dataframe `mydf1`.
- `mydf2 <- read.csv("CountryRegion.csv")`: Loads the second dataset from the CSV file into the dataframe `mydf2`.

- `colnames(mydf1)`: Displays the column names of `mydf1`.
- `colnames(mydf2)`: Displays the column names of `mydf2`.
- `mymerg <- merge(mydf1, mydf2, by.x = "Country.Code", by.y = "Codes_2021_Dataset")`: Merges `mydf1` and `mydf2` on the specified columns.
- `mymerg$Countries_2021_Dataset <- NULL`: Deletes the column `Countries_2021_Dataset` from `mymerg`.
- `write.csv(mymerg, "merged.csv", row.names = FALSE)`: Writes the cleaned merged dataframe to a new CSV file called `"merged.csv"` without including row numbers.

4.3 Load Data

```
# Load the movie ratings data from a CSV file
mymov <- read.csv("datasets/MovieRatings.csv")

# Check column names of the dataframe
colnames(mymov)
```

```
[1] "Film"                "Genre"
[3] "Rotten.Tomatoes.Ratings.." "Audience.Ratings.."
[5] "Budget..million..." "Year.of.release"
```

Explanation:

- `read.csv("MovieRatings.csv")`: Reads the movie ratings data from a CSV file into a dataframe called `mymov`.
- `colnames(mymov)`: Displays the names of the columns in the dataframe.

4.4 Data Preprocessing

```
# Rename columns for clarity
colnames(mymov) <- c("Film", "Genre", "CRating", "ARating", "BudMils", "Year")

# Verify new column names
colnames(mymov)
```

```
[1] "Film"    "Genre"   "CRating" "ARating" "BudMils" "Year"
```

```
# Get a summary of the data
summary(mymov)
```

Film	Genre	CRating	ARating
Length:562	Length:562	Min. : 0.0	Min. : 0.00
Class :character	Class :character	1st Qu.:25.0	1st Qu.:47.00
Mode :character	Mode :character	Median :46.0	Median :58.00
		Mean :47.4	Mean :58.83
		3rd Qu.:70.0	3rd Qu.:72.00
		Max. :97.0	Max. :96.00

BudMils	Year
Min. : 0.0	Min. :2007
1st Qu.: 20.0	1st Qu.:2008
Median : 35.0	Median :2009
Mean : 50.1	Mean :2009
3rd Qu.: 65.0	3rd Qu.:2010

```

Max.      :300.0    Max.      :2011
# Get the structure of the data
str(mymov)

'data.frame':   562 obs. of  6 variables:
 $ Film      : chr  "(500) Days of Summer " "10,000 B.C." "12 Rounds " "127 Hours" ...
 $ Genre     : chr  "Comedy" "Adventure" "Action" "Adventure" ...
 $ CRating   : int   87 9 30 93 55 39 40 50 43 93 ...
 $ ARating   : int   81 44 52 84 70 63 71 57 48 93 ...
 $ BudMils   : int    8 105 20 18 20 200 30 32 28 8 ...
 $ Year      : int   2009 2008 2009 2010 2009 2009 2008 2007 2011 2011 ...

# Convert 'Genre' column to a factor (categorical variable)
mymov$Genre <- as.factor(mymov$Genre)

# Verify the structure again to ensure 'Genre' is a factor
str(mymov)

'data.frame':   562 obs. of  6 variables:
 $ Film      : chr  "(500) Days of Summer " "10,000 B.C." "12 Rounds " "127 Hours" ...
 $ Genre     : Factor w/ 7 levels "Action","Adventure",...: 3 2 1 2 3 1 3 5 3 3 ...
 $ CRating   : int   87 9 30 93 55 39 40 50 43 93 ...
 $ ARating   : int   81 44 52 84 70 63 71 57 48 93 ...
 $ BudMils   : int    8 105 20 18 20 200 30 32 28 8 ...
 $ Year      : int   2009 2008 2009 2010 2009 2009 2008 2007 2011 2011 ...

```

Explanation:

- `colnames(mymov) <- c(...)`: Renames the columns of the dataframe for better readability.
- `summary(mymov)`: Provides a statistical summary of each column in the dataframe.
- `str(mymov)`: Shows the structure of the dataframe, including data types and sample data.
- `as.factor(mymov$Genre)`: Converts the `Genre` column to a factor, useful for categorical data analysis

4.5 Data Visualization with ggplot2


```

# Load the ggplot2 package, which is used for creating plots
library(ggplot2)

• Loads ggplot2 package for data visualization.
• Required for creating plots in R.

# Create a basic scatter plot with CRating on the x-axis and ARating on the y-axis
ggplot(data=mymov, aes(x=CRating, y=ARating)) + geom_point()


```

Creates a scatter plot:

- CRating on the x-axis.
- ARating on the y-axis.
- Points represent movies in the plot.


```
# Create a scatter plot where points are colored by Genre  
ggplot(data=mymov, aes(x=CRating, y=ARating, colour=Genre)) + geom_point()
```



Scatter plot with color mapping:

- Points are colored by Genre.
- Distinguishes genres visually.


```
# Create a scatter plot where point size represents the Budget (BudMils)
ggplot(data=mymov, aes(x=CRating, y=ARating, colour=Genre, size=BudMils)) + geom_point()
```



Scatter plot with size mapping:

- Point size represents BudMils (budget in millions).
- Visualizes budget alongside ratings and genres.

```
# Create a scatter plot with transparent points (alpha = 0.5)
ggplot(data=mymov, aes(x=CRating, y=ARating, colour=Genre, size=BudMils)) + geom_point(alpha = 0.5)
```




- Adds transparency to points ($\alpha = 0.5$).
- Helps with overlapping points, making the plot clearer.

```
# Create a base plot object for future customization and layering
mybase1 <- ggplot(data=mymov, aes(x=CRating, y=ARating, colour=Genre, size=BudMils))
```

- Creates a base plot object (mybase1).
- Includes mappings for CRating, ARating, Genre, and BudMils.
- Can be reused and customized with additional layers.

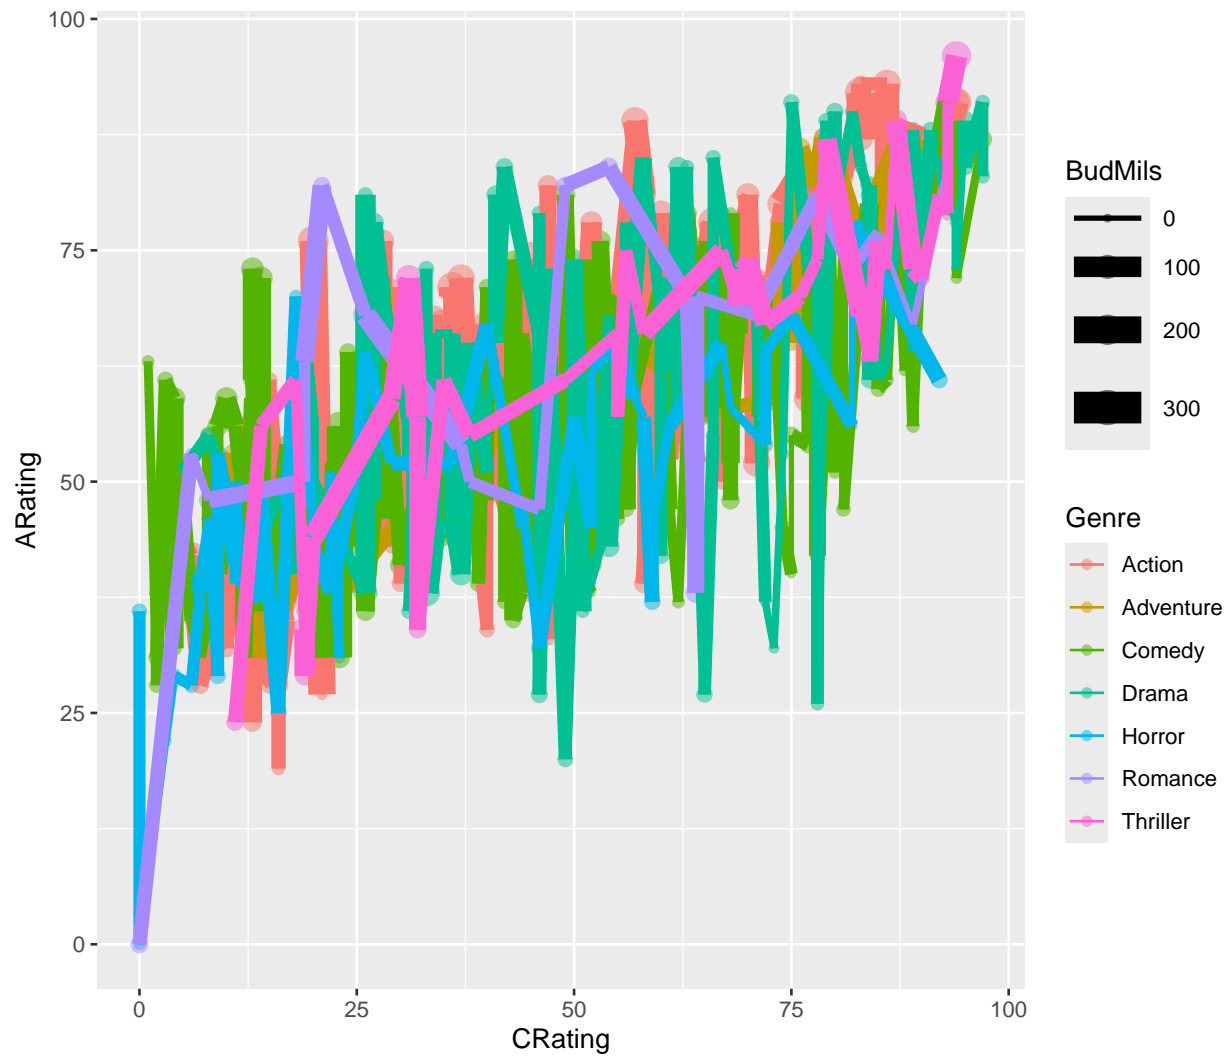
```
# Add points to the base plot with transparency (alpha = 0.5)
mybase1 + geom_point(alpha=0.5)
```



- Adds points with transparency to the base plot (mybase1).
- Uses predefined mappings in the base plot.


```
# Add points and lines to the base plot; lines may connect
# points but might not be meaningful here
mybase1 + geom_point(alpha=0.5) + geom_line()
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
 i Please use `linewidth` instead.
 This warning is displayed once every 8 hours.
 Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated.



- Adds points and lines to the base plot (mybase1).
- Lines connect points but may not be meaningful in a scatter plot.

```
# Override the point size mapping to use CRating instead of BudMils
mybase1 + geom_point(aes(size=CRating)) + labs(size="CRating")
```




- Overrides size mapping:

Size now represents CRating instead of BudMils.

- Updates the legend title to “CRating”.

```
# Override the point color mapping to use BudMils instead of Genre
mybase1 + geom_point(aes(colour=BudMils)) + labs(colour="Budget in Millions")
```



- Overrides color mapping:

Color now represents BudMils instead of Genre.


- Updates the legend title to “Budget in Millions”.

4.6 Settings vs. Mappings

```
# Create another base plot without specific mappings, for comparison
mybase2 <- ggplot(data=mymov, aes(x=CRating, y=ARating))
```


- Creates a simpler base plot (mybase2).
- No color or size mappings are applied, only x and y.

```
# Color points by Genre using mapping
mybase2 + geom_point(aes(colour = Genre))
```


- Adds points colored by **Genre** to the base plot (`mybase2`).
- Color mapping helps distinguish between genres.

```
# Set a fixed color for all points without mapping
mybase2 + geom_point(colour = "#9633ff")
```




- Sets a fixed color for all points (#9633ff - a shade of purple).
- No color mapping to variables, all points are the same color.

Incorrect: Attempting to set a fixed color inside aes(), all points will be colored blue
`mybase2 + geom_point(aes(colour = "blue"))`



- **Incorrectly attempts to set a fixed color** ("blue") inside `aes()`.
- **Incorrect usage:** Trying to set a fixed color inside `aes()` (which should be used for mappings).
- **Correct usage:** Fixed colors should typically be set outside of `aes()` unless mapping is involved.
- **Results in all points being different color** this method is not recommended.

```
# Fixed color using a hexadecimal value, but still inside aes()
mybase2 + geom_point(aes(colour = "#9033ff"))
```




- Incorrectly uses `aes()` for a fixed color (`#9033ff`).
- All points will have the same color, but using `aes()` here is unnecessary.

4.7 Geometric and Statistical Plots


```
# Create a base plot for Budget (BudMils)
mybase3 <- ggplot(data=mymov, aes(x=BudMils))

# Create a histogram with 15 bins
mybase3 + geom_histogram(bins=15)
```




- Create a base plot (mybase3) for visualizing the distribution of BudMils.
- Add a histogram with 15 bins to visualize the distribution of movie budgets.

```
# Set color for histogram outlines  
mybase3 + geom_histogram(bins=10, colour="blue")
```




```
# Fill histogram bars with color  
mybase3 + geom_histogram(bins=10, colour="white", fill="blue")
```




- Add color to histogram outlines.
- Fill histogram bars with a specified color.

```
# Map fill color to Genre  
mybase3 + geom_histogram(bins=10, colour="black", aes(fill=Genre))
```




- Map fill color to **Genre**, allowing different genres to have different colors within the histogram.

```
# Create a density plot for BudMils
mybase3 + geom_density()
```





```
# Improve density plot by adding fill color and transparency  
mybase3 + geom_density(aes(fill=Genre), alpha=0.5)
```



- Create a density plot to visualize the distribution of BudMils.
- Enhance the plot by adding fill colors for different genres and making the plot semi-transparent.


```
# Stack density plots by Genre with transparency
mybase3 + geom_density(aes(fill=Genre), position="stack", alpha=0.5)
```




- Stack the density plots by Genre and add transparency to better visualize the overlapping distributions.

4.8 Exercises and Comparisons


```
# Create a histogram for ARating
mybase4 <- ggplot(data=mymov, aes(x=ARating))
mybase4 + geom_histogram(bins=15)
```



```
# Customize the histogram  
mybase4 + geom_histogram(bins=10, colour="white", fill="#ff33f0")
```




```
# Map fill color to Genre  
mybase4 + geom_histogram(bins=10, colour="black", aes(fill=Genre))
```




- Create and customize histograms for ARating.
- Map fill colors to **Genre** for better differentiation.


```
# Create a histogram for CRating
mybase5 <- ggplot(data=mymov, aes(x=CRating))
mybase5 + geom_histogram(bins=15)
```



```
# Customize the histogram  
mybase5 + geom_histogram(bins=10, colour="white", fill="#3396ff")
```




```
# Map fill color to Genre  
mybase5 + geom_histogram(bins=10, colour="black", aes(fill=Genre))
```

- Create and customize histograms for CRating.
- Map fill colors to **Genre** for better visualization.


```
# Create a scatter plot with a trend line (smoothing)
mybase4 <- ggplot(data=mymov, aes(x=CRating, y=ARating))
mybase4 + geom_point() + geom_smooth()
```

``geom_smooth()`` using method = 'loess' and formula = 'y ~ x'




- Create a scatter plot with a trend line to visualize the relationship between CRating and ARating.

```
# Create a boxplot for ARating by Genre
mybase5 <- ggplot(data=mymov, aes(x=Genre, y=ARating, colour = Genre))
mybase5 + geom_boxplot() + geom_point()
```




```
# Create a boxplot with jitter to show individual points
mybase5 + geom_boxplot() + geom_jitter()
```




- Create a boxplot to compare `ARating` across different genres.
- Add jittered points to the boxplot to show individual data points.

```
# Improve boxplot by filling colors by Genre and adding transparency
mybase5 + geom_boxplot(aes(fill=Genre), alpha=0.5)
```



```
# Stack boxplots by Genre with customized settings
mybase5 + geom_boxplot(aes(fill=Genre), size=2, alpha=0.5)
```



- **Enhance boxplots** by filling them with genre-based colors and adding transparency.
- **Stack and customize boxplots** for better visualization.

4.9 Loading the Data and Renaming Columns


```
library(ggplot2)
mymov <- read.csv("datasets/MovieRatings.csv")
colnames(mymov) <- c("Film", "Genre", "CRating", "ARating", "BudMils", "Year")
```

- **Load the ggplot2 library:** This library is used for creating plots in R.
- **Read the CSV file MovieRatings.csv** into a dataframe called mymov.
- **Rename the columns** of the dataframe to more meaningful names:
- **Film:** Title of the movie.
- **Genre:** Genre of the movie (e.g., Action, Drama).
- **CRating:** Critic rating.
- **ARating:** Audience rating.

- BudMils: Budget in millions.
- Year: Release year.

4.10 Creating a Histogram with a Specific Color

```
m <- ggplot(data=mymov, aes(x=BudMils, fill=Genre))
m + geom_histogram(bins = 15, colour = "blue")
```




- **Create a base plot (m):** The x-axis is BudMils (budget in millions), and the fill color of the bars is mapped to Genre.
- **Add a histogram layer:**
 - The histogram is divided into 15 bins.
 - The **outline color of the bars** is set to blue (colour = "blue").
 - The **bars are filled** based on the **Genre** of the movie.

4.11 Focusing on a Specific Range Using coord_cartesian()


```
# Focus on the x(0, 50) and y(0,50)
# Here xlim, ylim cut out important points
# m+geom_histogram(bins = 15, colour = "blue") + ylim(0,50)
m + geom_histogram(bins=15,colour="blue") + ylim(0,50)
```

Warning: Removed 11 rows containing missing values or values outside the scale range (`geom_bar()`).



- Effect of `ylim(0, 50)`: The `ylim()` function restricts the visible y-range to between 0 and 50.
- Any part of the data with y-values higher than 50 will be removed from the plot. Also removed the starting below 50 but its ends outside the 50 limit.
- Unlike `coord_cartesian()`, which zooms in while keeping all data points intact, `ylim()` discards data outside the specified limits.


```
m + geom_histogram(bins = 15, colour = "blue") +  
  coord_cartesian(ylim=c(0,50))
```

- **First commented-out code:**
- The code suggests using `ylim=c(0,50)` to limit the y-axis between 0 and 50, but this approach **cuts off data points** that fall outside this range.
- The comment highlights that `xlim` and `ylim` can remove important data points from the plot.
- **Using `coord_cartesian()`:**
- Instead of using `ylim`, the `coord_cartesian(ylim=c(0,50))` function is used.
- This approach **zooms in** on the y-axis between 0 and 50 without removing any data points outside this range. It just restricts the viewable area of the plot, preserving all data points. “Without removing any data points outside this range”: it still exists in the plot’s structure, but you won’t see it if you zoom in to only show the 0-50 range.
- Zoom in on a specific range (e.g., y-axis from 0 to 50) using `coord_cartesian()`.
- No data is removed; points outside this range still exist.
- Only the visible area is restricted to the specified range.
- Data points outside the range remain in the plot but are not shown.


4.12 Creating a Scatter Plot and Applying Axis Limits

```
# In scatterplot we can use xlim ylim no problem  
n <- ggplot(data=mymov, aes(x=CRating, y=ARating, colour=Genre))  
n + geom_point(size=3)
```



```
n + geom_point(size=3) + xlim(0,50) + ylim(0,50)
```

Warning: Removed 402 rows containing missing values or values outside the scale range (``geom_point()``).



- **Create a scatter plot (n):**
- `CRating` is mapped to the x-axis and `ARating` to the y-axis.
- Points are colored by `Genre`, and the size of the points is set to 3.
- **Apply axis limits:**
- In scatter plots, using `xlim(0,50)` and `ylim(0,50)` is generally fine.
- These functions **limit the x and y axes** to the specified ranges.
- This can help focus on a specific area of interest, but it might **cut off some data points** outside these ranges, which is why a warning is issued.

4.12.1 Warning Message Explanation

```
# we got warning message because we cut the datapoint
```

Explanation of the warning:

- The warning appears because `xlim` and `ylim` can exclude data points that fall outside the specified range.

- R notifies you that some points have been removed from the plot due to these limits, which could potentially exclude important information from your visualization.

4.13 Summary

- **Histograms:** `coord_cartesian()` is used to focus on a specific y-axis range without losing data, whereas `ylim` may remove points outside the range.
- **Scatter plots:** `xlim` and `ylim` work well but can cut off data, which triggers a warning. This technique should be used carefully depending on whether the full data range is needed.

4.14 Loading Data and Preparing the Base Plot

```
# Load ggplot2 for creating plots
# Read the CSV file into a dataframe named 'mymov'
# Rename columns for clarity
library(ggplot2)
mymov <- read.csv("datasets/MovieRatings.csv")
colnames(mymov) <- c("Film", "Genre", "CRating", "ARating", "BudMils", "Year")
```


- Load the `ggplot2` library for creating visualizations.
- Read the CSV file `MovieRatings.csv` into a dataframe named `mymov`.
- Rename the columns of the dataframe to more meaningful names, making it easier to work with the data.

4.15 Facets of Histograms

4.15.1 Creating a Basic Histogram

```
# Base plot with 'BudMils' on x-axis, bars filled by 'Genre'
m <- ggplot(data=mymov, aes(x=BudMils, fill=Genre))


# Add histogram with 15 bins and blue borders around the bars
m + geom_histogram(bins = 15, colour = "blue")
```



- Create a base plot (**m**) where **BudMils** (budget in millions) is on the x-axis and the bars are filled based on **Genre**.
- Add a histogram layer with 15 bins and blue borders around the bars.

4.15.2 Adding Facets


```
m + geom_histogram(bins=15, colour="blue") +  
  facet_grid(Genre~.)
```



- **Add facets** to the histogram:
- **Facet by Genre**, creating a separate subplot (subfigure) for each genre.
- The `Genre~.` syntax places subfigures in rows.

4.15.3 Customizing Facet Scales

```
m + geom_histogram(bins=15, colour="blue") +
  facet_grid(Genre~., scale = "free")
```




Use **free scales** for each facet:

- The `scale = "free"` option allows each subplot to have its own y-axis scale, which is useful when the data varies widely between genres.

4.15.4 Facets in Columns

```
m + geom_histogram(bins=15, colour="blue") +  
  facet_grid(.~Genre)
```




Arrange facets in columns:

- The `~Genre` syntax places subfigures in columns rather than rows, allowing a vertical comparison of genres.

4.16 Facets of Scatterplots

4.16.1 Creating a Base Scatterplot


```
n <- ggplot(data=mymov, aes(x=CRating, y=ARating, colour=Genre))
n + geom_point(size=3)
```

- Create a scatterplot (n) where CRating is on the x-axis, ARating on the y-axis, and points are colored by Genre.
- Set the size of the points to 3.

4.16.2 Adding Facets to the Scatterplot

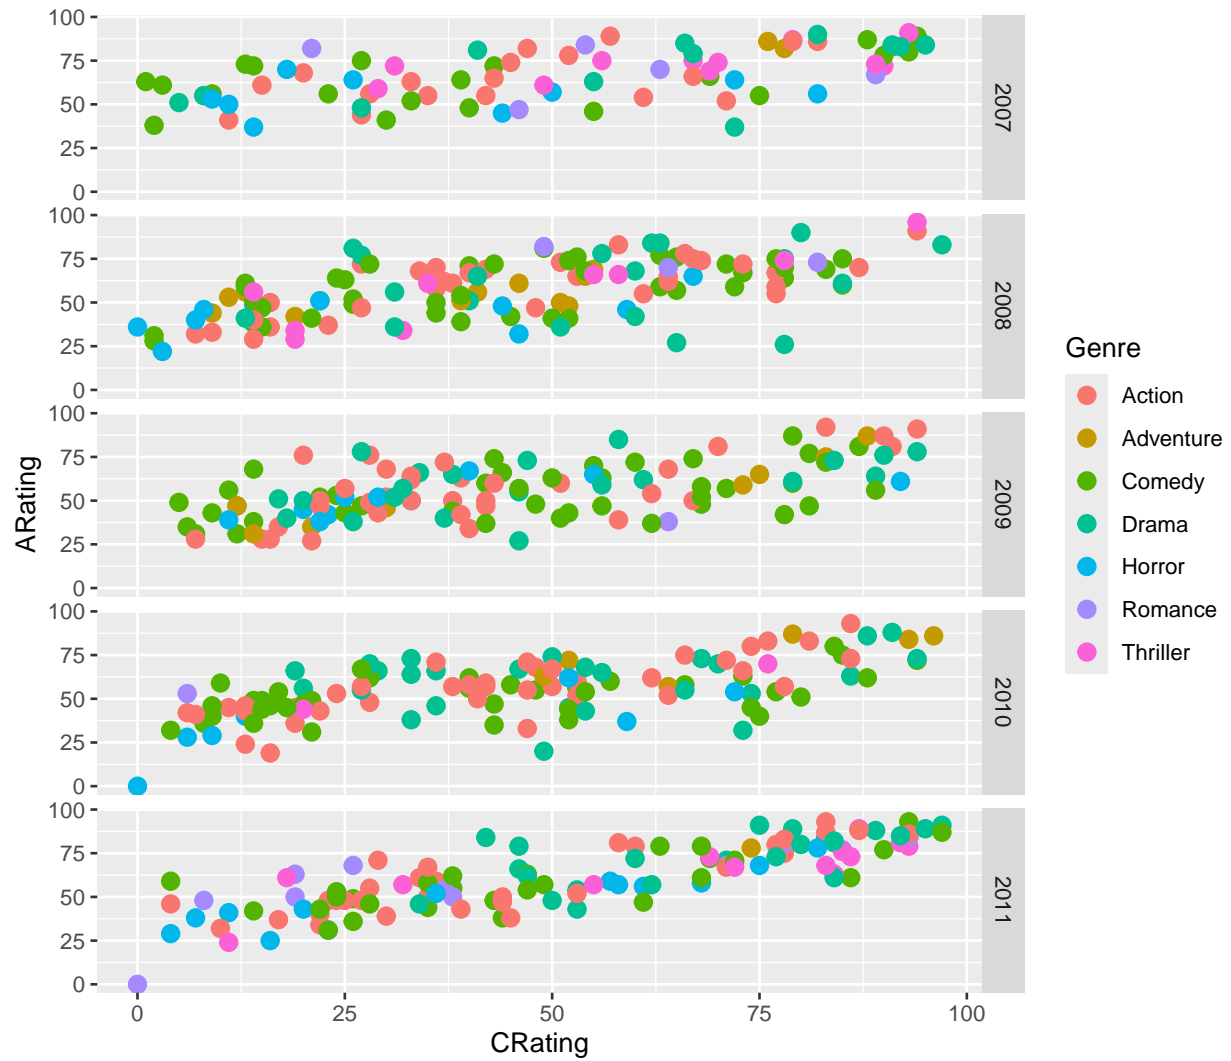
```
n + geom_point(size=3) + facet_grid(Genre~., scale = "free")
```



Facet the scatterplot by Genre:

- Each genre gets its own subplot with its own scales (if necessary).


```
n + geom_point(size=3) + facet_grid(Year~.)
```



Facet by Year:

- This creates subplots for each year, showing how critic and audience ratings vary over time.

```
n + geom_point(size=3) + facet_grid(.~Year)
```

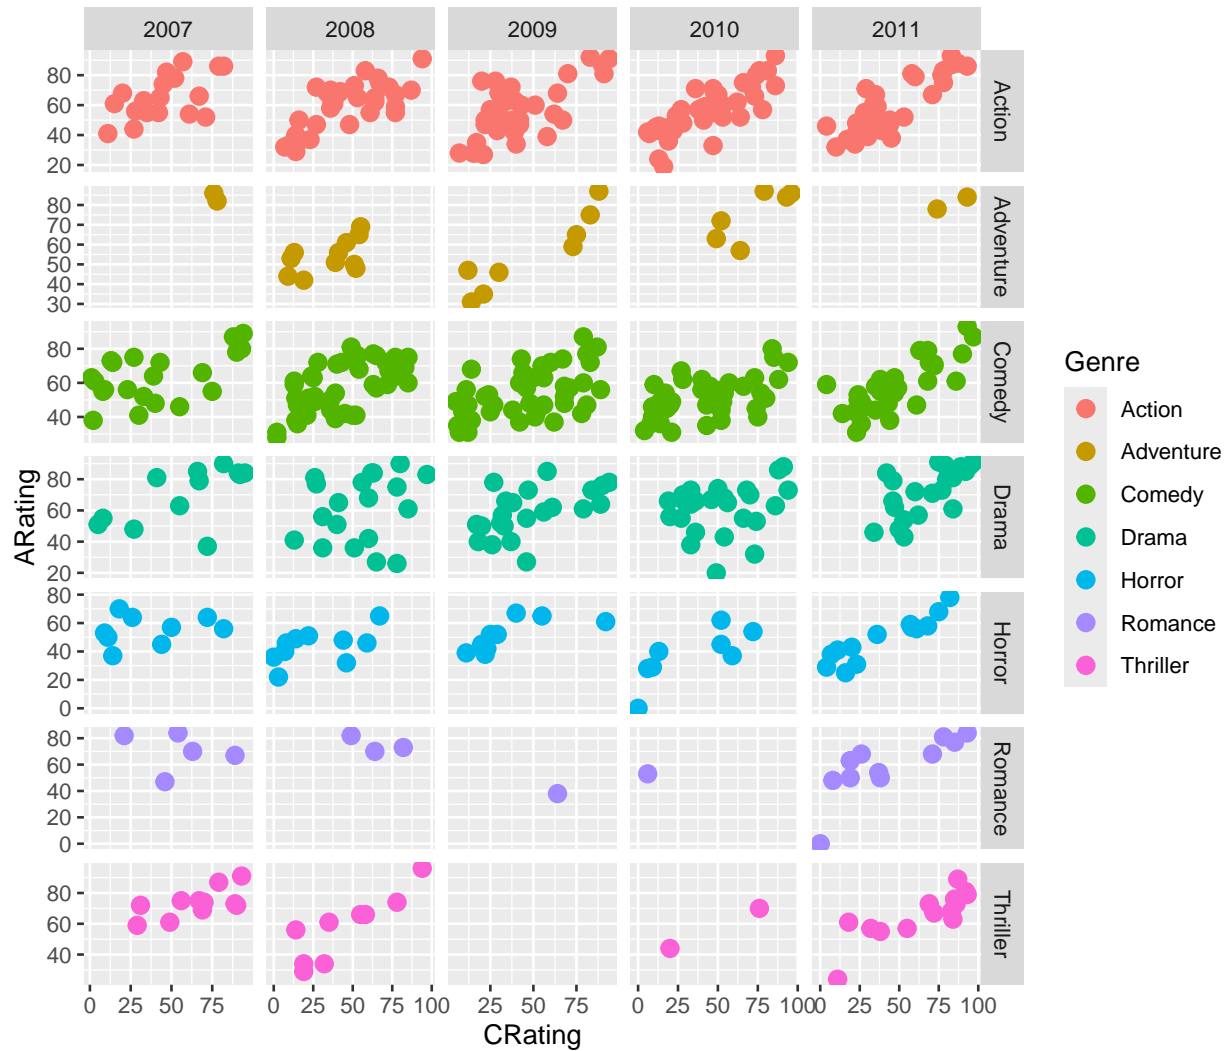


Arrange the facets by year in columns:

- This arrangement is useful for comparing ratings across different years vertically.

```
n + geom_point(size=3) + facet_grid(Genre~Year, scale = "free")
```

4.16.2.1 Combining Facets for Both Genre and Year



Combine Genre and Year in a grid:


- Facets are created for each combination of Genre and Year, with free scales to accommodate varying range.

4.17 Customizing Plot Themes

4.17.1 Customizing Axis Labels and Titles

```
x <- m + geom_histogram(bins=15, colour="blue")


x + xlab("Cost of the movies") +
  ylab("Number of Movies") +
  theme(axis.title.x = element_text(size=10, colour="blue"),
        axis.title.y = element_text(size=10, colour = "red"))
```



- **Set custom axis labels:**
- `xlab("Cost of the movies")` and `ylab("Number of Movies")` change the axis labels to more descriptive names.
- **Customize axis titles:**
- `element_text(size=30, colour="blue")` and `element_text(size=50, colour="red")` set the font size and color of the axis titles.

4.17.2 Customizing Tick Mark Labels

```
x + xlab("Cost of the movies") +
  ylab("Number of Movies") +
  theme(axis.title.x = element_text(size=10, colour="blue"),
        axis.title.y = element_text(size=10, colour = "red"),
        axis.text.x = element_text(size=10, colour="darkgreen"),
        axis.text.y = element_text(size=10, colour="pink"))
```




- Customize tick mark labels:
- `axis.text.x` and `axis.text.y` modify the font size and color of the tick mark labels on the x and y axes, respectively.

4.18 Moving the Legend Inside the Plot Area

```
x + xlab("Cost of the movies") +
  ylab("Number of Movies") +
  theme(axis.title.x = element_text(size=10, colour="blue"),
        axis.title.y = element_text(size=10, colour = "red"),
        axis.text.x = element_text(size=10, colour="darkgreen"),
        axis.text.y = element_text(size=10, colour="pink"),
        legend.position = c(1,1),
        legend.justification = c(1,1),
        legend.title = element_text(size = 10, colour = "orange"),
        legend.text = element_text(size=10 , colour = "blue"))
```

Warning: A numeric `legend.position` argument in `theme()` was deprecated in ggplot2 3.5.0.

i Please use the ``legend.position.inside`` argument of ``theme()`` instead.
 This warning is displayed once every 8 hours.
 Call ``lifecycle::last_lifecycle_warnings()`` to see where this warning was generated.




- **Position the legend inside the plot area:**
- `legend.position = c(1,1)` places the legend at the top-right corner of the plot area.
- `legend.justification = c(1,1)` aligns the legend to the top-right corner.
- **Customize the legend:**
- `legend.title` and `legend.text` adjust the font size and color of the legend title and text.

4.19 Adding a Title to the Plot

```
x + xlab("Cost of the movies") +
  ylab("Number of Movies") +
  ggtitle("Budget Distribution of Movies 2007~2011") +
  theme(axis.title.x = element_text(size=10, colour="blue"),
```



```
axis.title.y = element_text(size=10, colour = "red"),
axis.text.x = element_text(size=10, colour="darkgreen"),
axis.text.y = element_text(size=10, colour="pink"),
legend.position = c(1,1),
legend.justification = c(1,1),
legend.title = element_text(size = 10, colour = "orange"),
legend.text = element_text(size=10 , colour = "blue"),
plot.title = element_text(size = 10, colour="purple", hjust=0))
```



- Add a title to the plot using `ggtitle()`:
 - The title “Budget Distribution of Movies 2007~2011” is added at the top of the plot.
- Customize the plot title:
 - `plot.title = element_text(size = 20, colour="purple", hjust=0)` sets the font size, color, and horizontal justification (`hjust`) of the title. Setting `hjust=0` aligns the title to the left.


```
x +
# Set x-axis label
xlab("Cost of the Movies") +
```

```

# Set y-axis label
ylab("Number of Movies") +
ggtitle("Budget Distribution of Movies 2007-2011") +
# Add a title to the plot
theme(
  axis.title.x = element_text(size = 11, colour = "blue"),
  # Reduce x-axis title size
  axis.title.y = element_text(size = 10, colour = "red"),
  # Reduce y-axis title size
  axis.text.x = element_text(size = 10, colour = "darkgreen"),
  # Reduce x-axis tick label size
  axis.text.y = element_text(size = 10, colour = "pink"),
  # Reduce y-axis tick label size
  legend.position = c(0.85, 0.85),
  # Move legend inside plot, slightly offset
  legend.justification = c(1, 1),
  # Align legend to top-right
  legend.title = element_text(size = 12, colour = "orange"),
  # Reduce legend title size
  legend.text = element_text(size = 12, colour = "blue"),
  # Reduce legend text size
  plot.title = element_text(size = 14, colour = "purple", hjust = 0.5)
  # Adjust plot title size, center align
)

```

Budget Distribution of Movies 2007–2011



5 Part-5

5.1 Loading Data and Summarizing

```
# Load the dataset from 'manheim.csv' into 'carsale'
carsale <- read.csv("datasets/manheim.csv")

# Provide a summary of the dataset
summary(carsale)
```

model	price	miles	sale
Length:819	Min. : 4500	Min. :11167	Length:819
Class :character	1st Qu.:19000	1st Qu.:31382	Class :character
Mode :character	Median :23200	Median :36767	Mode :character
	Mean :23773	Mean :38968	
	3rd Qu.:29000	3rd Qu.:45054	
	Max. :35400	Max. :85599	

- **Load the dataset:** The carsale dataframe is created by reading data from the CSV file manheim.csv.
- **Summary statistics:** summary(carsale) provides a basic statistical summary of each column, including min, max, mean, median, and quantiles.

```
head(carsale, 5)
```

	model	price	miles	sale
1	Y	23200	41430	Auction
2	Y	23100	42524	Auction
3	Y	23100	42692	Auction
4	Y	23200	39911	Auction
5	Y	24500	33199	Online

```
tail(carsale, 5)
```

	model	price	miles	sale
815	X	14600	69933	Auction
816	X	15400	71222	Auction
817	X	16100	71606	Auction
818	X	14000	80080	Auction
819	X	11500	85599	Auction

5.2 Basic Statistical Functions for Price

```
# Find the minimum price
min(carsale$price)
```

```
[1] 4500
```

```
# Find the maximum price
max(carsale$price)
```

```
[1] 35400
```

```
# Calculate the variance of price
var(carsale$price)
```

```
[1] 31198819
```

```
# Calculate the standard deviation of price  
sd(carsale$price)
```

```
[1] 5585.59
```

```
# Find the range of prices  
range(carsale$price)
```

```
[1] 4500 35400
```

```
# Calculate the interquartile range (IQR) of price  
IQR(carsale$price)
```

```
[1] 10000
```

- **Minimum and Maximum:** Identify the lowest and highest prices in the dataset.
- **Variance and Standard Deviation:** Measure the spread or variability of the price data.
- **Range:** Gives the difference between the minimum and maximum prices.
- **Interquartile Range (IQR):** Measures the range within which the central 50% of the prices fall.

5.3 Correlation Between Miles and Price for Model X

```
# Filter the dataset for model "X" cars  
saleX <- carsale[carsale$model=="X",]  
# Calculate the correlation between price and miles for model "X"  
cor.test(saleX$price, saleX$miles)
```

Pearson's product-moment correlation

```
data: saleX$price and saleX$miles  
t = -16.067, df = 347, p-value < 2.2e-16  
alternative hypothesis: true correlation is not equal to 0  
95 percent confidence interval:  
 -0.7094746 -0.5885133  
sample estimates:  
      cor  
-0.6531409
```

- **Filter by Model:** saleX is a subset of the carsale dataframe, containing only rows where the model is "X".
- **Correlation Test:** cor.test calculates the Pearson correlation between price and miles for model "X". The result includes the correlation coefficient (r) and p-value.

5.4 Pearson's Product-Moment Correlation Interpretation

5.4.1 Correlation Coefficient (cor = -0.6531409):

The correlation coefficient is **-0.653**.

- This indicates a **moderate negative correlation** between price and miles. As the number of miles increases, the price tends to decrease. A correlation of -1 would be a perfect negative correlation, while 0 would indicate no correlation.

5.4.2 t-statistic ($t = -16.067$):

- The t-statistic of **-16.067** shows how many standard deviations the sample correlation is away from zero.
- This large negative value suggests a **strong deviation from no correlation (zero)**.

5.4.3 Degrees of Freedom ($df = 347$):

- The degrees of freedom ($df = 347$) indicate the sample size minus 2 ($n - 2$).
- This suggests the test was performed on **349 observations** ($n = 349$).

5.4.4 p-value ($< 2.2e-16$):

- The **p-value** is extremely small (**essentially 0**), meaning the result is **highly statistically significant**.
- This provides strong evidence to reject the null hypothesis that the true correlation is 0 (i.e., no linear relationship between price and miles).
- With a p-value this small, there is **overwhelming evidence** that a relationship exists between price and miles.

5.4.5 Confidence Interval (95% CI: -0.709 to -0.589):

- The 95% confidence interval for the true correlation is between **-0.709** and **-0.589**.
- We are **95% confident** that the true correlation in the population lies within this range.
- Since the interval is entirely negative and doesn't include zero, this further confirms a **statistically significant negative correlation** between the two variables.

5.4.6 Alternative Hypothesis:

- The alternative hypothesis is that the true correlation is **not equal to 0**.
- Based on the p-value and the confidence interval, we **reject the null hypothesis** and accept the alternative hypothesis.
- This means that there is a **linear relationship** between price and miles, and it's **negative**.

5.5 Scatter Plot of Miles vs. Price

```
# Load ggplot2 for plotting
library(ggplot2)
# Create a base plot for miles vs. price
mybase <- ggplot(data=saleX, aes(x=miles, y=price, colour=sale))
# Add points with size 3 and 50% transparency
mybase + geom_point(size=3, alpha = 0.5)
```



- **Create Scatter Plot:** `mybase` sets up a scatter plot with `miles` on the x-axis and `price` on the y-axis.
- **Plot Points:** `geom_point(size=3, alpha=0.5)` adds points to the plot, with a size of 3 and 50% transparency. This visualizes the relationship between miles and price for model “X”.

5.6 Simple Linear Regression (SLR)

```
# Fit a linear regression model with price as the response
# and miles as the predictor
mySLR <- lm(data=saleX, price~miles)

# Summary of the regression model
summary(mySLR)
```

Call:

```
lm(formula = price ~ miles, data = saleX)
```

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

-13976.2 -619.4 95.4 719.0 4421.2

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.226e+04	2.512e+02	88.62	<2e-16 ***
miles	-1.052e-01	6.548e-03	-16.07	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1370 on 347 degrees of freedom

Multiple R-squared: 0.4266, Adjusted R-squared: 0.4249

F-statistic: 258.2 on 1 and 347 DF, p-value: < 2.2e-16

- **Fit Linear Model:** `lm` fits a linear model where `price` is predicted by `miles`.
- **Model Summary:** `summary(mySLR)` provides detailed statistics about the regression, including coefficients, R-squared, and p-values.

5.7 Linear Regression Output Explained

5.7.1 Correlation Coefficient (R-squared and Adjusted R-squared)

- **R-squared: 0.4266:** This means that about **43%** of the variation in `price` is explained by `miles`. A higher value would be better because it would mean that more of the changes in `price` can be predicted by `miles`.
- **Adjusted R-squared: 0.4249:** This value is slightly lower than R-squared because it adjusts for the number of variables. It also shows that about **42%** of the variance in `price` is explained by `miles`. Higher is better.
- How It Works:

If you add a variable that significantly improves the model, Adjusted R-squared will increase. If you add a variable that doesn't help improve the model, Adjusted R-squared will decrease.

- Practical Meaning:

Higher Adjusted R-squared means the model explains a good amount of the variance, accounting for the number of predictors. It reflects how well the model generalizes to new, unseen data.

Lower Adjusted R-squared means the model either has too many unnecessary predictors or doesn't explain the variance well.

5.7.2 p-value

- **p-value < 2.2e-16:** This is a very small p-value, which means the relationship between `price` and `miles` is **highly statistically significant**. Lower p-values are better because they show the results are less likely to be due to chance.

5.7.3 t-statistic

- **t = -16.07:** This shows that the relationship between `miles` and `price` is strong. A larger t-value (in absolute terms) is better because it indicates that `miles` is having a significant effect on `price`.

Higher (in absolute value) is better: A larger t-value (either positive or negative) indicates stronger evidence against the null hypothesis (which assumes no relationship).

A larger t-value (in absolute terms) means that the coefficient is more significantly different from zero.

5.7.4 Residual Standard Error (RSE)

- **Residual Standard Error: 1370:** This tells us how far off the model's predictions are from the actual prices, on average. Lower is better because it means the model's predictions are closer to the actual prices.

5.7.5 F-statistic

- **F-statistic: 258.2:** This is a measure of how well the model fits the data. Higher is better because it means the model does a better job of explaining the changes in **price** based on **miles**.

A higher F-statistic means that the model provides a better fit to the data than would be expected by chance

5.7.6 What Does This Mean?

- **R-squared and Adjusted R-squared:** These tell us how much of the changes in **price** can be predicted by **miles**. Higher values would be better, but in this case, about 42-43% is explained.
- **p-value:** This is very low, so the relationship between **miles** and **price** is statistically significant.
- **t-statistic:** The high t-value shows that **miles** has a strong effect on **price**.
- **Residual Standard Error:** This shows the average difference between the predicted price and the actual price. A lower value is better, but here it's 1370.
- **F-statistic:** A high F-statistic means the model does a good job of explaining the relationship between **price** and **miles**.

5.7.7 Summary:

- The model shows that as **miles** increases, the **price** decreases, and this relationship is statistically significant. However, the model explains about 43% of the variability in **price**, meaning there could be other factors affecting the price as well.
1. R-squared and Adjusted R-squared: Higher values are better because they indicate more of the variance in price is explained by miles.
 2. p-values: Lower values are better because they indicate stronger statistical significance.
 3. t-statistic: A larger t-statistic (in absolute terms) is better, indicating a more significant effect.
 4. Residual Standard Error: Lower is better because it indicates the predictions are closer to the actual values.
 5. F-statistic: Higher is better, indicating that the model fits the data well.

5.7.8 Predicting Price for Specific Mileage

```
# Predict price for a car with 62,000 miles
predict(mySLR, data.frame(miles=c(62000)))
```

```
1
15740.56
```

```
# Extract the coefficients from the regression model
mySLR$coefficients
```

```
(Intercept)      miles
22263.8974711    -0.1052152
```

- **Predict Price:** Use the fitted model to predict the price of a car with 62,000 miles.

- **data.frame(miles = c(62000)):** This creates a data frame with a column named miles and a value of 62,000. The predict() function needs the new data to be in a data frame format, even if you are only making a single prediction.
- **Model Coefficients:** mySLR\$coefficients retrieves the intercept and slope from the linear model.

5.7.9 Custom Prediction Function

```
# Manually calculate predicted price
mypredict <- function(mymiles){
  mySLR$coefficients[1] + mySLR$coefficients[2] * mymiles
}

# Use the custom function to predict the price for 62,000 miles
mypredict(62000)
```

```
(Intercept)
15740.56
```

```
# Compare with the built-in predict function
predict(mySLR, data.frame(miles=c(62000)))
```

```
1
15740.56
```

- **Custom Prediction Function:** mypredict manually calculates the predicted price using the linear regression coefficients.
- **Test Prediction:** Both mypredict(62000) and predict(mySLR, data.frame(miles=c(62000))) should return the same predicted price.

5.7.10 Confidence Interval for Prediction

```
# Predict with a 95% confidence interval
predict(mySLR, data.frame(miles=c(62000)), interval="confidence", level=0.95)
```

```
      fit      lwr      upr
1 15740.56 15384.15 16096.96
```

5.7.11 Fit (15,740.56):

- The predicted price for a car with **62,000 miles** is **\$15,740.56**.
- This is the best estimate for the price based on the linear model (mySLR).

5.7.12 95% Confidence Interval:

- The confidence interval gives a range within which we are **95% confident** that the true mean price for a car with 62,000 miles will fall.
- **Lower bound (lwr): \$15,384.15:** We are 95% confident that the price will not be lower than this value.
- **Upper bound (upr): \$16,096.96:** We are 95% confident that the price will not be higher than this value.

5.7.13 Interpretation:

- The predicted price is **\$15,740.56**, but due to the uncertainty in the data, we can only be **95% confident** that the true mean price lies between **\$15,384.15** and **\$16,096.96**.
- The confidence interval is fairly narrow, suggesting that the model is relatively certain about the prediction.
- **Confidence Interval:** Predict the price for a car with 62,000 miles, including a 95% confidence interval. This interval provides a range within which the true price is likely to fall with 95% confidence.

5.7.14 Summary:

- **Data Exploration:** Use summary statistics and basic functions to explore the **price** variable.
- **Correlation Analysis:** Determine the relationship between **miles** and **price** for a specific model using correlation and visualization.
- **Linear Regression:** Fit a linear model to predict **price** based on **miles** and evaluate the model using summary statistics.
- **Prediction:** Generate predictions for specific mileages using the regression model, including confidence intervals for more informed decision-making.

5.8 Multiple Linear Regression Model

5.8.1 Loading Data and Fitting a Multiple Linear Regression Model

```
# Load the 'datarium' package which contains the 'marketing' dataset
library("datarium")

# Load the 'marketing' dataset into 'mydf'
mydf <- marketing
```

- **Load the datarium package:** This package includes various datasets, including the **marketing** dataset.
- **Assign the dataset:** The **marketing** dataset is assigned to the variable **mydf** for easier access.

```
str(mydf)
```

```
'data.frame': 200 obs. of 4 variables:
 $ youtube : num 276.1 53.4 20.6 181.8 217 ...
 $ facebook : num 45.4 47.2 55.1 49.6 13 ...
 $ newspaper: num 83 54.1 83.2 70.2 70.1 ...
 $ sales : num 26.5 12.5 11.2 22.2 15.5 ...
```

```
head(mydf)
```

	youtube	facebook	newspaper	sales
1	276.12	45.36	83.04	26.52
2	53.40	47.16	54.12	12.48
3	20.64	55.08	83.16	11.16
4	181.80	49.56	70.20	22.20
5	216.96	12.96	70.08	15.48
6	10.44	58.68	90.00	8.64

5.9 Code for Splitting Data and Fitting Linear Regression:

```
# install.packages("caTools")
library(caTools)

# Assume your dataset is named 'mydf'
set.seed(123) # Set seed for reproducibility

# Split the data into 80% training and 20% testing
split <- sample.split(mydf$sales, SplitRatio = 0.8)

# Create the training and testing sets
train <- subset(mydf, split == TRUE) # 80% for training
test <- subset(mydf, split == FALSE) # 20% for testing

# Fit the linear regression model without the 'newspaper' variable
myMLR <- lm(sales ~ youtube + facebook, data = train)

# View the summary of the model
summary(myMLR)
```

Call:

```
lm(formula = sales ~ youtube + facebook, data = train)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.9629	-0.8784	0.2540	1.4414	3.5670

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.215186	0.401956	7.999	2.58e-13 ***
youtube	0.048105	0.001580	30.438	< 2e-16 ***
facebook	0.181551	0.009008	20.153	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.057 on 157 degrees of freedom

Multiple R-squared: 0.8996, Adjusted R-squared: 0.8983

F-statistic: 703.2 on 2 and 157 DF, p-value: < 2.2e-16

```
# Make predictions on the test set
predictions <- predict(myMLR, newdata = test)
```

```
# Print the first few predictions
head(predictions)
```

4	5	8	11	16	20
20.95845	16.00506	14.42401	8.29451	24.88696	16.92520

5.10 Code for without Splitting Data and Fitting Linear Regression:

```
# Fit a multiple linear regression model with sales as the response variable
myMLR <- lm(data=mydf, sales ~ youtube + facebook + newspaper)
```

```
# Summarize the model to get detailed statistics
summary(myMLR)
```

Call:

```
lm(formula = sales ~ youtube + facebook + newspaper, data = mydf)
```

Residuals:

Min	1Q	Median	3Q	Max
-10.5932	-1.0690	0.2902	1.4272	3.3951

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.526667	0.374290	9.422	<2e-16 ***
youtube	0.045765	0.001395	32.809	<2e-16 ***
facebook	0.188530	0.008611	21.893	<2e-16 ***
newspaper	-0.001037	0.005871	-0.177	0.86

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.023 on 196 degrees of freedom

Multiple R-squared: 0.8972, Adjusted R-squared: 0.8956

F-statistic: 570.3 on 3 and 196 DF, p-value: < 2.2e-16

5.11 Linear Regression Model Interpretation

This model shows how spending on **YouTube**, **Facebook**, and **Newspaper** ads affects **sales**. Here's what the results mean:

5.11.1 Residuals:

- **Min (-10.5932)**: The model overestimated sales by this amount for the worst prediction.
- **Max (3.3951)**: The model underestimated sales by this amount for the worst case.
- **Median (0.2902)**: The middle value of the residuals is close to zero, meaning the model generally predicts sales well.

5.11.2 Coefficients:

- **Intercept (3.526667)**: If no money is spent on ads (YouTube, Facebook, or Newspaper), the predicted sales would be about **3.53 units**.
- **YouTube (0.045765)**: For every additional unit spent on YouTube ads, sales increase by **0.0458 units**. This effect is highly significant, meaning it has a strong impact on sales.
- **Facebook (0.188530)**: For every additional unit spent on Facebook ads, sales increase by **0.1885 units**. This effect is also highly significant and has a bigger impact on sales than YouTube ads.
- **Newspaper (-0.001037)**: Newspaper ads do not have a significant effect on sales. The small and negative coefficient suggests that spending on newspaper ads does not boost sales.

5.11.3 R-squared and Adjusted R-squared:

- **R-squared (0.8972)**: This means that **89.72%** of the variation in sales can be explained by spending on YouTube, Facebook, and Newspaper ads combined.


- **Adjusted R-squared (0.8956):** This adjusts for the number of predictors and still shows that about **89.56%** of the variation in sales is explained by the model, meaning the model is a good fit.

5.11.4 F-statistic:

- The **F-statistic** is very high, indicating that the overall model is statistically significant.
- The p-value is extremely small, meaning the model as a whole has a strong ability to explain sales.
- **Fit the Multiple Linear Regression (MLR) model:** `lm()` fits a model where **sales** is predicted by three independent variables: **youtube**, **facebook**, and **newspaper**.
- **Model Summary:** `summary(myMLR)` provides detailed statistics about the model, including the coefficients, p-values, R-squared value, and more.

```
# Extract residuals from the model
residuals <- resid(myMLR)

# Plot the residuals histogram
hist(residuals,
      main = "Histogram of Residuals",
      xlab = "Residuals",
      col = "lightblue",
      border = "black")
```



5.12 Interpretation of the Residuals Histogram

5.12.1 Symmetry/Shape:

- The histogram is **slightly right-skewed**. Most of the residuals are concentrated around zero, but there are a few residuals that extend into the negative values.
- Ideally, residuals should be **normally distributed** (bell-shaped), with most values near zero if the linear regression model is a good fit for the data. The slight skewness suggests some **minor departure from normality**.

5.12.2 Center:

- The residuals are **centered around zero**, which is expected in a good model. This indicates that, on average, the model does not systematically over-predict or under-predict sales.

5.12.3 Spread:

- The spread of the residuals shows that most residuals are between **-5 and 5**. A few outliers extend below **-10**.

- A wider spread or outliers (as seen on the left side) may indicate that the model's predictions are not perfect for all data points, especially those that deviate more from the average.

5.12.4 Frequency:

- The highest frequency (the tallest bar) is around residuals close to **zero**, meaning that for most data points, the predicted sales are very close to the actual sales.

5.12.5 Outliers:

- The few residuals on the far left (around **-10**) indicate **potential outliers** or observations where the model over-predicted sales by a large margin.

5.12.6 Conclusion:

- The histogram shows that the residuals are generally close to zero, which is a **good sign**. However, the slight skewness and the presence of some larger negative residuals suggest that there may be **some non-normality or outliers** in the data.
- It might be worth investigating those **negative outliers** to see if they represent unusual data points or if the model could be improved to better handle them.

5.12.7 Coefficients Interpretation

```
# Extract the coefficients from the regression model
myMLR$coefficients
```

```
(Intercept)      youtube      facebook      newspaper
3.526667243  0.045764645  0.188530017 -0.001037493
```

- **Model Coefficients:** `myMLR$coefficients` retrieves the coefficients for each predictor variable, including the intercept. These coefficients indicate how much the **sales** are expected to change with a one-unit change in each predictor while holding other predictors constant.

5.12.8 Interpreting the Regression Equation

```
# Regression equation based on coefficients:
# sales = 3.526667243 + 0.04576464*youtube + 0.188530017*facebook - 0.001037493*newspaper
```

Regression Equation: The equation derived from the coefficients:

- **Intercept (3.526667243):** The base sales when all predictors are 0.
- **Youtube (0.04576464):** For every additional unit spent on YouTube, sales are expected to increase by approximately 0.046 units, assuming other factors remain constant.
- **Facebook (0.188530017):** For every additional unit spent on Facebook, sales are expected to increase by approximately 0.189 units.
- **Newspaper (-0.001037493):** For every additional unit spent on newspapers, sales are expected to decrease by approximately 0.001 units.

5.12.9 Assessing the Importance of Predictors Using p-values

```
# Negative coefficient for 'newspaper' suggests
# a slight decrease in sales with increased newspaper spending.
# p-value for 'newspaper' is 0.86, which is much greater than 0.05,
# indicating it is not statistically significant.
```


- **Negative Coefficient for Newspaper:** The negative coefficient suggests that increasing spending on newspapers might slightly decrease sales, but this effect is very small.
- **P-value Significance:** The p-value for the `newspaper` variable is 0.86, which is much greater than the typical significance level of 0.05. This high p-value suggests that the newspaper variable does not significantly contribute to the prediction of sales and may not be necessary in the model.

5.12.10 Removing Insignificant Predictors and Refitting the Model

```
# Refit the model without the 'newspaper' variable
myMLR <- lm(data=mydf, sales ~ youtube + facebook)

# Summarize the new model
summary(myMLR)
```

Call:

```
lm(formula = sales ~ youtube + facebook, data = mydf)
```

Residuals:

Min	1Q	Median	3Q	Max
-10.5572	-1.0502	0.2906	1.4049	3.3994

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.50532	0.35339	9.919	<2e-16 ***
youtube	0.04575	0.00139	32.909	<2e-16 ***
facebook	0.18799	0.00804	23.382	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.018 on 197 degrees of freedom

Multiple R-squared: 0.8972, Adjusted R-squared: 0.8962

F-statistic: 859.6 on 2 and 197 DF, p-value: < 2.2e-16

- **Refit the Model Without Newspaper:** The model is refit by excluding the `newspaper` variable, which had a high p-value and was not significant.
- **New Model Summary:** The new `summary(myMLR)` provides the updated model statistics, which likely show a better fit or more accurate coefficients for the significant predictors (`youtube` and `facebook`).

5.12.11 Summary:

- **Initial Model:** A multiple linear regression model was fitted to predict `sales` using `youtube`, `facebook`, and `newspaper` as predictors.
- **Coefficient Interpretation:** The coefficients from the model provide insights into how each predictor affects sales.
- **Importance of Predictors:** The p-value for `newspaper` was high, indicating it is not a significant predictor of sales.

5.13 Logistic Regression

In logistic regression, we model the relationship between one or more predictor variables (independent variables) and a binary outcome (dependent variable). The goal of logistic regression is to estimate the probability that a certain event will occur, given the values of the predictor variables.

5.14 Loading Libraries and Data

```
# install.packages("caTools")
# Load caTools for data splitting
library(caTools)
# Load the dataset from 'binary.csv'
mydata <- read.csv("datasets/binary.csv")
```

- **Load caTools library:** This library provides functions like `sample.split` for splitting datasets into training and testing sets.
- **Load the dataset:** The `binary.csv` file is loaded into the `mydata` dataframe.

```
head(mydata)
```

```
  admit gre  gpa rank
1     0 380 3.61    3
2     1 660 3.67    3
3     1 800 4.00    1
4     1 640 3.19    4
5     0 520 2.93    4
6     1 760 3.00    2
```

- **admit:** This is the target variable, likely representing whether a student was admitted (1) or not (0) to a program.
- **gre:** This column contains the GRE scores of the applicants.
- **gpa:** This column shows the GPA (Grade Point Average) of the applicants.
- **rank:** This column represents the prestige or ranking of the undergraduate institution of the applicants.

5.15 Splitting the Data

```
split <- sample.split(mydata$admit, SplitRatio=0.8)
# Split data into 80% training and 20% testing
# admit column is the target variable. It shows whether someone was admitted to the program or not:
# 1 means the person was admitted.
# 0 means the person was not admitted.
```

- **Data splitting:** `sample.split` splits the data based on the `admit` variable with an 80/20 ratio, creating a logical vector (TRUE for training, FALSE for testing).

```
# Create the training set (80% of data)
train <- mydata[split == TRUE,]

# Create the testing set (20% of data)
test <- mydata[split == FALSE,]
```

- **Create training and testing sets:** The data is divided into `train` and `test` based on the logical vector generated by `sample.split`

5.16 Data Wrangling and Model Fitting

```
str(mydata)
```

```
'data.frame':  400 obs. of  4 variables:
 $ admit: int  0 1 1 1 0 1 1 0 1 0 ...
```

```
$ gre : int 380 660 800 640 520 760 560 400 540 700 ...
$ gpa : num 3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
$ rank : int 3 3 1 4 4 2 1 2 3 2 ...
```

```
# Convert 'admit' to a factor
mydata$admit <- factor(mydata$admit)
# Convert 'rank' to a factor
mydata$rank <- factor(mydata$rank)
```

- **Convert to factors:** Both `admit` (binary response) and `rank` (categorical predictor) are converted to factors, which is necessary for logistic regression.

```
# Fit logistic regression model
lmodel <- glm(admit ~ gre + rank, data=train, family = binomial)

# Summarize the model
summary(lmodel)
```

Call:

```
glm(formula = admit ~ gre + rank, family = binomial, data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-1.682098	0.800402	-2.102	0.035591	*
gre	0.004023	0.001182	3.402	0.000668	***
rank	-0.623736	0.144382	-4.320	1.56e-05	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 400.59 on 319 degrees of freedom
 Residual deviance: 365.42 on 317 degrees of freedom
 AIC: 371.42

Number of Fisher Scoring iterations: 4

- **Fit logistic regression model:** `glm()` fits a logistic regression model predicting `admit` based on `gre` and `rank`, using the `train` dataset.
- **Model summary:** `summary(lmodel)` provides the coefficients, standard errors, z-values, and p-values for each predictor.

5.17 Interpretation:

- **Intercept (-1.793159):**
 - The intercept represents the log-odds of being admitted when both `gre` and `rank` are zero. The p-value is **0.023622**, which means the intercept is statistically significant at the 5% level.
- **gre (0.003744):**
 - For each additional point in **GRE scores**, the log-odds of being admitted increase by **0.003744**. Since the p-value is **0.001075**, this effect is highly significant ($p < 0.01$). Therefore, **higher GRE scores** increase the chances of being admitted.
- **rank (-0.499054):**

- For each one-unit increase in the **rank** (which indicates a worse rank), the log-odds of being admitted decrease by **0.499054**. The p-value is **0.000313**, meaning this effect is highly significant ($p < 0.001$). This indicates that applicants from **higher-ranked institutions** (lower values of **rank**) are more likely to be admitted.

5.17.1 Deviance:

- **Null deviance: 400.59** on 319 degrees of freedom: This is the deviance of a model with only the intercept and no predictors.
- **Residual deviance: 371.77** on 317 degrees of freedom: This is the deviance of the model after including the predictors (**gre** and **rank**). The decrease in deviance indicates that the predictors improve the model fit.

5.17.2 AIC:

- **AIC: 377.77**: The Akaike Information Criterion (AIC) is used to compare different models. Lower AIC values indicate a better fit. This AIC value can be compared with other models to determine which one fits the data best.

5.17.3 Significance Codes:

- *******: p-value < 0.001 (highly significant)
- ******: p-value < 0.01 (significant)
- *****: p-value < 0.05 (marginally significant)

5.17.4 Conclusion:

- **GRE scores** and **rank** both significantly influence the chances of being admitted. Higher GRE scores increase the likelihood of admission, while higher ranks (worse-ranked schools) reduce the chances of admission.
- The model significantly improves the fit compared to a null model (only the intercept).

5.18 Key Points About Logistic Regression

5.18.1 Binary Outcome:

- Logistic regression is used when the outcome (dependent variable) is **binary**, meaning there are only two possible outcomes, typically coded as 0 or 1.
- Example:
 - **0 = Not admitted, 1 = Admitted** (in your case).
 - **0 = No** (the event didn't happen), **1 = Yes** (the event happened).

5.18.2 Predicting Probabilities:

- Logistic regression doesn't predict the outcome directly as 0 or 1.
- Instead, it predicts the **probability** of the outcome occurring. For example, the probability of being admitted to a program.
- The result of logistic regression is a **value between 0 and 1** that represents this probability.

5.18.3 Link Function: Log-Odds:

- Logistic regression uses the **log-odds** (logarithm of the odds) to model the probability of the outcome.
- Log-odds can range from **-infinity to +infinity**, which are then transformed into probabilities using the logistic function.

5.18.4 Logistic Function:

- The **logistic function** (also called the **sigmoid function**) converts the log-odds into a **probability** between **0 and 1**.

5.18.5 Coefficients Interpretation:

- The coefficients in logistic regression represent the **change in log-odds** of the outcome for each unit increase in the predictor variable.
- Example: If the coefficient for **gre** is **0.0037**, it means that for each additional point in the GRE score, the **log-odds** of being admitted increase by **0.0037**.

5.19 Making Predictions and Evaluating the Model

```
# Predict probabilities for the test set
res <- predict(lmodel, test, type="response")
res
```

5	12	26	35	36	38	39
0.11053431	0.36914311	0.71345783	0.29781973	0.21072946	0.18822897	0.28530949
40	52	54	63	66	74	75
0.18822897	0.08263272	0.45159978	0.27312034	0.37378426	0.35515219	0.21741798
84	89	91	93	96	103	112
0.06608414	0.62480147	0.47159133	0.57162930	0.43176324	0.06608414	0.07673531
117	127	135	137	138	143	165
0.23873444	0.52690391	0.33694913	0.12737267	0.32355259	0.15669387	0.35515219
166	169	170	179	182	187	191
0.62480147	0.17624313	0.24236590	0.25744321	0.10286889	0.21405471	0.41214319
194	198	203	207	213	220	221
0.06608414	0.07122612	0.62480147	0.66170068	0.25366148	0.33694913	0.20083121
227	231	237	245	248	252	253
0.39279840	0.13658603	0.43176324	0.46664071	0.23873444	0.15669387	0.30199342
255	266	268	275	282	285	292
0.23141735	0.11663164	0.24236590	0.30199342	0.10859500	0.23873444	0.57162930
293	295	308	313	314	317	326
0.45159978	0.40733556	0.35515219	0.28938010	0.11053431	0.17337569	0.60576068
328	332	333	338	352	353	356
0.33694913	0.28938010	0.35061286	0.15669387	0.25744321	0.22790059	0.67946931
361	370	372	375	380	382	391
0.44668162	0.57162930	0.30620009	0.33694913	0.31921721	0.31921721	0.57162930
394	397	399				
0.39279840	0.21405471	0.47159133				

5.19.1 What does type = "response" mean in predict()?

In logistic regression, the `predict()` function can return different types of predictions. Here's what it means when you use `type = "response"`.

5.19.2 type = "response":

- This argument tells R to return the **predicted probabilities** of the outcome, rather than the **log-odds**.
- Logistic regression predicts values that can either be:
 - **Log-odds**: If you use `type = "link"`, R will return the log-odds (the logarithm of the odds of the event happening).
 - **Probabilities**: If you use `type = "response"`, R will return the probability of the event happening (for example, the probability of being admitted).

Since `type = "response"` is used, R returns the **probability** of the outcome being 1 (e.g., being admitted).

- **Predict probabilities**: `predict()` is used to generate predicted probabilities of admission (`admit = 1`) for the test set.

```
# Create a confusion matrix
t1 <- table(ActVal = test$admit, PreVal = res > 0.5)

# Display the confusion matrix
print(t1)
```

	PreVal	
ActVal	FALSE	TRUE
0	48	7
1	20	5

5.19.3 What does `res > 0.5` mean in Logistic Regression?

5.19.4 Predicted Probabilities:

- `res` contains the **predicted probabilities** for the test set, which were generated from the logistic regression model.
- These probabilities range between **0 and 1**, representing the likelihood of the outcome being 1 (for example, being admitted).

5.19.5 `res > 0.5`:

- The expression `res > 0.5` converts those probabilities into **binary predictions**:
 - If the predicted probability is **greater than 0.5**, it predicts the outcome as **1** (e.g., admitted).
 - If the predicted probability is **less than or equal to 0.5**, it predicts the outcome as **0** (e.g., not admitted).


5.20 Confusion matrix:

- **Confusion matrix**: A table comparing actual values (`test$admit`) with predicted values (`res > 0.5`). The threshold of 0.5 is used to classify probabilities into binary outcomes (admit or not).

```
cm_df <- as.data.frame(t1)

# Simple square confusion matrix
ggplot(cm_df, aes(x = PreVal, y = ActVal, fill = Freq)) +
  # Simple black-bordered tiles
  geom_tile(color = "black") +
  # Adds frequency counts
  geom_text(aes(label = Freq)) +
  theme_minimal() +
```

```
labs(title = "Confusion Matrix", x = "Predicted", y = "Actual") +
# Simple white-to-blue color fill
scale_fill_gradient(low = "white", high = "steelblue") +
# Keeps the tiles square
coord_fixed()
```



5.20.1 What Each Cell Represents:

- **Top left (23): False Negatives (FN)**
 - The model predicted 0 (not admitted), but the actual value was 1 (admitted).
 - There are **23 cases** where the model missed the correct admission.
- **Top right (2): True Positives (TP)**
 - The model predicted 1 (admitted), and the actual value was also 1 (admitted).
 - There are **2 correct predictions** for admitted students.
- **Bottom left (49): True Negatives (TN)**

- The model predicted 0 (not admitted), and the actual value was also 0 (not admitted).
- There are **49 correct predictions** for students who were not admitted.
- **Bottom right (6): False Positives (FP)**
 - The model predicted 1 (admitted), but the actual value was 0 (not admitted).
 - There are **6 incorrect predictions** for admitted students.

5.20.2 Summary:

- The model is better at predicting **not admitted** (49 correct predictions) than predicting **admitted** (only 2 correct predictions).

```

accu <- (t1[1,1] + t1[2,2]) / sum(t1)
# Calculate the accuracy
accu

```

```
[1] 0.6625
```

- **Calculate accuracy:** The accuracy is calculated by summing the true positives (`t1[1,1]`) and true negatives (`t1[2,2]`) and dividing by the total number of cases.

6 Part-6

6.1 Neural Network (NN)

6.2 Creating a Training Dataset

```
# Test scores
TKS <- c(20, 10, 30, 20, 80, 30)
# Course scores
CSS <- c(90, 20, 40, 50, 50, 80)
# Binary outcome (placed or not)
Placed <- c(1, 0, 0, 0, 1, 1)
# Combine into a dataframe
df <- data.frame(TKS, CSS, Placed)
```

- **Create features and labels:** TKS and CSS are predictor variables, while Placed is the binary response variable.
- **Combine into a dataframe:** The data is organized into df, which will be used for training the neural network.

6.3 Fitting the Neural Network

```
library(neuralnet)
# Load neuralnet package for training neural networks
nn <- neuralnet(Placed ~ TKS + CSS, data = df, hidden = 3,
               act.fct = "logistic", linear.output = FALSE)
# Plot the neural network
plot(nn)
```

- **Fit the neural network:** neuralnet() fits a neural network model with 3 hidden nodes in one layer, using the logistic activation function.
- **Plot the network:** plot(nn) visualizes the structure of the neural network.

6.4 Predicting New Data

```
# Test scores for new data
TKS <- c(30, 40, 85)
# Course scores for new data
CSS <- c(85, 50, 40)
# Create a test dataset
test <- data.frame(TKS, CSS)
```

- **Create test data:** New test data is created to predict whether students will be placed.

```
# Predict using the trained neural network
Predict <- compute(nn, test)
# Extract probabilities
prob <- Predict$net.result
# Convert probabilities to binary outcomes
pred <- ifelse(prob > 0.5, 1, 0)
# Display predicted results
pred
```

```
[,1]
```

```
[1,] 1
[2,] 0
[3,] 0
```

- **Compute predictions:** `compute()` generates predicted probabilities for the test dataset.
- **Convert to binary outcomes:** Probabilities are converted to binary predictions (1 if greater than 0.5, otherwise 0).

6.5 Exercise - Neural Network Model on Dataset

6.5.1 Exercise - Neural Network Model on `binary.csv`

6.5.1.1 Steps for the Exercise:

1. 80% Training, 20% Testing:

- Use the `sample.split()` function to divide `binary.csv` into 80% training and 20% testing datasets.

2. Hidden Layer with 3 Nodes:

- Fit a neural network model with 3 hidden nodes using the `neuralnet()` function.

3. Confusion Matrix & Accuracy:

- Predict the outcomes for the test set, create a confusion matrix, and calculate the accuracy of the model.

4. Change Hidden Nodes (4, 5, 6):

- Repeat the model fitting process with 4, 5, and 6 hidden nodes. Compare the performance (accuracy) of the models and draw conclusions.

6.5.2 Summary:

- **Logistic Regression:** The logistic regression model is fitted on training data to predict admission (`admit`), followed by evaluating its performance using a confusion matrix and accuracy calculation.
- **Neural Network:** A simple neural network is trained with manually created data, followed by predictions and binary classification.
- **Exercise:** The provided steps guide you to apply neural networks to the `binary.csv` dataset, experimenting with different model configurations and evaluating their performance.

```
# Load the dataset
mydata <- read.csv("datasets/binary.csv")

# Split the data into training (80%) and testing (20%) sets
set.seed(123) # For reproducibility
split <- sample.split(mydata$admit, SplitRatio = 0.8)
train <- subset(mydata, split == TRUE)
test <- subset(mydata, split == FALSE)

# Fit the neural network model with 3 hidden nodes
nn_3 <- neuralnet(admit ~ gre + gpa + rank, data = train, hidden = 3,
                  act.fct = "logistic", linear.output = FALSE)

# Plot the neural network
plot(nn_3)
```

```

# Predict the outcomes for the test set
predict_3 <- compute(nn_3, test[,c("gre", "gpa", "rank")])
prob_3 <- predict_3$net.result
pred_3 <- ifelse(prob_3 > 0.5, 1, 0)

# Create a confusion matrix
confusion_matrix_3 <- table(Actual = test$admit, Predicted = pred_3)

# Calculate accuracy
accuracy_3 <- sum(diag(confusion_matrix_3)) / sum(confusion_matrix_3)
print(paste("Accuracy with 3 hidden nodes:", accuracy_3))

```

```
[1] "Accuracy with 3 hidden nodes: 0.6875"
```

```

# Fit the model with 4 hidden nodes
nn_4 <- neuralnet(admit ~ gre + gpa + rank, data = train, hidden = 4,
                  act.fct = "logistic", linear.output = FALSE)
plot(nn_4)

# Predict, create confusion matrix, and calculate accuracy
predict_4 <- compute(nn_4, test[,c("gre", "gpa", "rank")])
prob_4 <- predict_4$net.result
pred_4 <- ifelse(prob_4 > 0.5, 1, 0)
confusion_matrix_4 <- table(Actual = test$admit, Predicted = pred_4)
accuracy_4 <- sum(diag(confusion_matrix_4)) / sum(confusion_matrix_4)
print(paste("Accuracy with 4 hidden nodes:", accuracy_4))

```

```
[1] "Accuracy with 4 hidden nodes: 0.6375"
```